

Affordable Real Style Transfer: Training Spectral Flow Matching on an RTX 3060 in Minutes

Anonymous submission

Abstract

We present Spectral Flow Matching (SFM), a style transfer method that trains in 3 minutes on a single 12 GB RTX 3060. It reaches CLIP-S = 0.7213 and LPIPS = 0.2868 on 750 pairs across five WikiArt domains, dual-beating a 436-minute Mamba baseline and a commercial diffusion API. Consumer-grade style transfer typically produces a color filter, not real brushwork: SaMam reports CLIP-S = 0.5816, below the identity baseline 0.6933. We trace this failure to a *degeneration attractor* in Euclidean latent space, where content and style gradients share a basis. A multiplicative attenuation collapses the style signal. The attractor dissolves once transport moves to the orthogonal wavelet domain. SFM applies a Haar wavelet to the SDXL VAE latent and locks the LL subband as a content anchor. It runs flow matching only on LH, HL, HH and injects style once at the trajectory endpoint via per-subband whitening-and-coloring. The resulting model has 903K parameters, trains in 3 minutes, and transfers real brushstroke direction and edge texture rather than a global color shift.

1 Introduction

Real style transfer on consumer-grade hardware is hard. Modern diffusion-based methods such as CSGO train on 210K paired images with A100 clusters, which is out of reach for academic replication. Methods that do train on consumer hardware, such as CUT and SaMam, are slow: CUT takes 322 minutes and SaMam takes 436 minutes on an RTX 3090. Worse, on the 750-pair protocol (which includes identity pairs), SaMam’s reported CLIP-S of 0.5816 falls below the identity baseline of 0.6933. This means it performs content reconstruction, not style transfer. Its low LPIPS is therefore a style-failure symptom, not a content-preservation win. In practice, consumer-grade outputs are typically a global color shift: the content image is preserved, but brushwork, texture, and high-frequency style cues are absent. When the model is pushed harder for style fidelity, the content tears apart. Practitioners respond by adding modules — cross-attention, style gates, FiLM, optimal-transport adapters, discriminators — and by enlarging compute. The color-filter failure nonetheless persists on unpaired domains.

The Pareto deadlock. On the standard CLIP-S (style fidelity) vs. LPIPS (content preservation) plane, unsupervised latent methods cluster on a single one-dimensional frontier.

The empirical trade-off ratio is about 1:8: every +0.01 CLIP-S costs ≈ 0.08 LPIPS. We call this the *Pareto deadlock*: no amount of module stacking or hyperparameter tuning within the Euclidean latent space moves the operating point off this line. The frontier is not a capacity limit; it is a structural limit imposed by content–style entanglement in pixel/latent space.

The degeneration attractor. We trace the deadlock to a *degeneration attractor*: a parameter manifold on which the learned velocity field becomes style-independent and small in norm. Three multiplicative mechanisms keep the model on this manifold. These are a collapsed style gate (learned scalar $\rightarrow 0.05$), a uniform cross-attention (entropy near maximum), and a GroupNorm that washes out per-channel style statistics. Their product is ≈ 0.016 , matching the observed residual style coefficient. Adding more modules does not help because each new module is itself pulled to the attractor.

The wavelet escape. Our central insight is that the attractor exists because content and style share the same Euclidean basis. Once we apply a 2D Haar wavelet transform, the LL subband carries global structure and tone, while LH/HL/HH carry directional brushwork. Content preservation becomes a constraint on LL only, and style transfer becomes a problem on LH/HL/HH only. The two objectives no longer compete. The loss landscape, previously a rugged valley, becomes nearly convex. The model devotes 100% of its capacity to learning brushwork, and convergence takes minutes rather than hours.

Spectral Flow Matching. We propose *Spectral Flow Matching* (SFM). SFM applies a Haar wavelet transform to the SDXL VAE latent and locks the LL subband as a content anchor. It learns a flow-matching velocity field only on the LH/HL/HH subbands and injects style once at the trajectory endpoint by per-subband whitening-and-coloring. The base manifold (LL) is invariant; style transfer happens entirely along the high-frequency fibers.

Results. SFM has 903K parameters and trains in 3 minutes on a single 12 GB RTX 3060. On a 750-pair protocol across five WikiArt domains it reaches CLIP-S = 0.7213 and LPIPS = 0.2868. It dual-beats SaMam by +0.14 CLIP-S at 1/141 the training cost, and beats the commercial Seedream

4.5 API by +0.19 LPIPS. The transfer is real brushwork rather than a color filter. Brushstroke direction, edge sharpness, and texture granularity move with the style while the content image’s color palette is preserved.

Contributions.

- We diagnose the Pareto deadlock in unsupervised latent style transfer and trace it to a multiplicative degeneration attractor (Section 3.2).
- We propose *Spectral Flow Matching* (SFM), a flow-matching velocity field on Haar wavelet subbands. SFM uses (i) *base locking* that keeps the LL subband fixed, (ii) *fiber flow* that injects style only into LH/HL/HH via per-subband whitening-and-coloring, and (iii) *end-of-trajectory injection* that decouples ODE integration from style blending (Section 3).
- On a 750-pair protocol across five WikiArt domains, SFM reaches CLIP-S = 0.7213 and LPIPS = 0.2868 with 903K parameters trained in 3 minutes on a single RTX 3060. It dual-beats SaMam (141× slower) and the commercial Seedream 4.5 API (Section 4).

2 Related Work

Neural style transfer. Optimization-based methods (Gatys *et al.* 2015, 2016) match Gram matrices per-image. AdaIN (Huang and Belongie 2017) and WCT (Li *et al.* 2017) align channel statistics in a single forward pass. SANet (Park and Lee 2019) and StyTr2 (Deng *et al.* 2022) add attention. These methods run in real time but produce color-filter outputs on domains with strong brushwork (e.g., Impressionism, Ukiyo-e). The cause is that they operate in Euclidean feature space, where content and style are entangled.

Diffusion-based style transfer. StyleID (Chung *et al.* 2024) injects style into a frozen diffusion model without training; SDEdit (Meng *et al.* 2021) adds noise and denoises. Both achieve high CLIP-S but damage content (LPIPS > 0.45). CSGO uses 210K paired images and A100 clusters, which is economically out of reach for academic replication. SaMam (Liu and others 2024b) and SaMST (Liu and others 2024a) train Mamba-based style encoders. On the 750-pair protocol (which includes identity pairs), SaMam’s CLIP-S (0.5816) falls below the identity baseline (0.6933). This indicates that it performs content reconstruction, not style transfer.

Flow matching and Schrödinger bridges. Flow matching (Lipman *et al.* 2022) learns a velocity field that transports one distribution to another via an ODE, avoiding simulation. Schrödinger bridges add an entropy regularizer and a stochastic term. We use the deterministic $\sigma=0$ special case (rectified flow) and condition the velocity field on a Haar wavelet decomposition of the latent. Prior wavelet-based generative work (Phung *et al.* 2023) uses wavelets for multi-scale sampling; we use them for content–style *disentanglement*, which to our knowledge is new.

Wavelet methods in vision. The Haar wavelet (Daubechies 1992) is the simplest orthogonal wavelet and has been used in image compression, denoising, and multi-scale architectures. In generative modeling, wavelet decompositions enable coarse-to-fine sampling and reduce sampling cost. Our use is structurally different: we exploit the orthogonality of the Haar basis to *decouple* two competing loss terms (content preservation and style injection) into orthogonal subspaces. This turns a non-convex reconciliation problem into two independent convex ones. The change is a coordinate-system change, not a multi-scale trick.

Disentanglement and coordinate-system methods. Content–style disentanglement is a long-standing goal. AdaIN and WCT disentangle via channel statistics; CycleGAN and CyCADA (Hoffman *et al.* 2018) use cycle consistency; contrastive methods (CUT (Park *et al.* 2020)) use patch-level negative pairs. These methods all operate in the Euclidean feature space, where content and style gradients compete. Our approach is orthogonal to this line: we do not add a new disentanglement *loss*. We change the *coordinate system* so that the existing losses act on disjoint subspaces. The wavelet transform provides this coordinate change for free, with no learned parameters and no additional loss terms.

Mamba-based style transfer. SaMam (Liu and others 2024b) and SaMST (Liu and others 2024a) adapt the Mamba state-space model (Gu and Dao 2023) to style encoding. On the 750-pair protocol (which includes identity pairs), SaMam’s reported CLIP-S (0.5816) falls below the identity baseline (0.6933). This indicates content reconstruction rather than style transfer. Its low LPIPS (0.2434) is therefore a style-failure symptom, not a content-preservation win. Our method uses a lightweight per-domain token bank instead of a Mamba encoder, and trains 141× faster.

Optimal transport and statistics matching. AdaIN matches diagonal covariance; WCT matches full covariance via whitening-and-coloring. We apply WCT per-subband at the trajectory endpoint only. This avoids the iterative residual decay $r_n = (1-\alpha)^n$ (Proposition 4) that makes per-step injection equivalent across α values.

3 Method

3.1 Setup

Let $x_c \in \mathbb{R}^{H \times W \times 3}$ be a content image and $s \in \{1, \dots, K\}$ a target style domain. We learn a mapping $G_\theta(x_c, s) = \mathcal{D}(\hat{z}_s)$ where $\mathcal{D}$ is a frozen VAE decoder and $\hat{z}_s$ is a stylized latent. We work in the latent space of a pre-trained SDXL VAE: $z_c = \mathcal{E}(x_c) \in \mathbb{R}^{4 \times 32 \times 32}$. The transport is a deterministic ODE (rectified flow) from $z_0 = z_c$ to $z_1 \sim p_{\text{style}}$.

Why latent space? We operate in the SDXL VAE latent space rather than pixel space for two reasons. First, the $4 \times 32 \times 32$ latent reduces training cost by roughly $192\times$ compared to 512×512 pixels, making consumer-grade training feasible. Second, the VAE’s compressed representation provides a more semantically meaningful feature




Figure 1: Spectral Flow Matching. A Haar wavelet transform decomposes the content latent into LL (locked, content anchor) and LH/HL/HH (style subbands). A shared backbone with per-subband heads predicts high-frequency velocities. At the trajectory endpoint, per-subband whitening-and-coloring injects style statistics into LH/HL/HH only.

space for style. Pixel-space features are dominated by low-level texture and color statistics that entangle with content. VAE latents abstract away low-level detail and expose style-relevant semantic structure. Latent diffusion produces more coherent synthesis than pixel-space diffusion at equal compute (Rombach *et al.* 2022).

Flow matching. We learn a velocity field $v_\theta(z_t, t, s)$ with $t \in [0, 1]$. The training target is the rectilinear interpolation

$$z_t = (1 - t)z_0 + tz_1, \quad u_t = z_1 - z_0, \quad (1)$$

and the flow-matching loss is

$$\mathcal{L}_{\text{FM}}(\theta) = \mathbb{E}_{t \sim \mathcal{U}(0,1)} \|v_\theta(z_t, t, s) - u_t\|_2^2. \quad (2)$$

At inference, we integrate the ODE with an 8-step Heun solver (second-order Runge-Kutta):

$$z_{t+\Delta t} = z_t + \frac{\Delta t}{2} [v_\theta(z_t, t, s) + v_\theta(z_t + \Delta t v_\theta(z_t, t, s), t + \Delta t, s)]. \quad (3)$$

Heun gives a structural degree of freedom that Euler does not. Switching Euler $\rightarrow$ Heun improves CLIP-S by +0.0056 at fixed α ; switching Heun $\rightarrow$ RK4 adds nothing (saturation). The solver order is a Pareto-breaking knob; schedule shape, α , and capacity are not.

3.2 Why Euclidean Fails

We diagnose why the Euclidean approach fails before presenting our spectral solution. Consider a velocity field $v_\theta(z, t, s)$ trained with loss $\mathcal{L} = w_{\text{FM}}\mathcal{L}_{\text{FM}} + w_{\text{SWD}}\mathcal{L}_{\text{SWD}}$ where $\mathcal{L}_{\text{SWD}}$ is a sliced-Wasserstein style term.

Definition 1 (Degeneration manifold). *The degeneration manifold is the set of parameters on which the velocity field is style-independent:*

$$\mathcal{M}_0 = \{\theta : v_\theta(z, t, s) = \bar{v}_\theta(z, t) = \mathbb{E}_{s \sim p(s)} [v_\theta(z, t, s)] \quad \forall s, z, t\}. \quad (4)$$

Equivalently, $\nabla_s v_\theta(z, t, s)|_{\theta \in \mathcal{M}_0} = 0$: the Jacobian of v_θ with respect to the style variable vanishes identically.

Theorem 1 ($\mathcal{M}_0$ is a local minimum under multiplicative attenuation). *Assume the loss decomposes as $\mathcal{L} = w_{\text{FM}}\mathcal{L}_{\text{FM}} + w_{\text{SWD}}\mathcal{L}_{\text{SWD}}$ with $w_{\text{FM}} \gg w_{\text{SWD}}$, and that three multiplicative attenuation factors act on the style gradient: a scalar gate $g(\theta) \in [0, 1]$, an attention entropy factor $h(\theta) \in [0, 1]$ (with $h = 1$ uniform), and a normalization washout factor $n(\theta) \in [0, 1]$. Let $\kappa(\theta) = g(\theta)h(\theta)n(\theta)$. If at $\theta_0 \in \mathcal{M}_0$:*

- (i) $\kappa(\theta_0) \leq \sqrt{w_{\text{SWD}}/w_{\text{FM}}}$ (style gradient too small to overcome content gradient),
- (ii) the Hessian $\nabla_{\theta^\perp}^2 \mathcal{L}_{\text{FM}}|_{\theta_0}$ is positive definite on the normal bundle $T_{\theta_0}^\perp \mathcal{M}_0$ (content loss is locally strictly convex in the escape direction),
- (iii) $\nabla^2 \mathcal{L}_{\text{SWD}}|_{\theta_0}$ has operator norm $\leq \kappa(\theta_0) \|\nabla^2 \mathcal{L}_{\text{FM}}\|$ (style loss is flat near $\mathcal{M}_0$),

then θ_0 is a strict local minimum of $\mathcal{L}$. The basin of attraction has radius $\Theta(\kappa(\theta_0))$.

Proof sketch. By (i), $\|\nabla_\theta \mathcal{L}_{\text{SWD}}\| \leq \kappa \|\nabla_\theta \mathcal{L}_{\text{FM}}\|$ at θ_0 , so the total gradient $\nabla \mathcal{L} = w_{\text{FM}}\nabla \mathcal{L}_{\text{FM}} + w_{\text{SWD}}\nabla \mathcal{L}_{\text{SWD}}$ is dominated by $w_{\text{FM}}\nabla \mathcal{L}_{\text{FM}}$, which vanishes on $\mathcal{M}_0$ by definition. By (ii), the second variation $\delta^2 \mathcal{L}_{\text{FM}} > 0$ on $T^\perp \mathcal{M}_0$; by (iii), $\delta^2 \mathcal{L}_{\text{SWD}}$ is a bounded perturbation that does not flip the sign of $\delta^2 \mathcal{L}$ within radius $O(\kappa)$. Hence $\delta^2 \mathcal{L} > 0$ in a neighborhood of θ_0 on $T^\perp \mathcal{M}_0$, making θ_0 a strict local min. Full proof in the supplement. $\square$ $\square$

In our Euclidean baseline, the three factors are measured empirically: gate $g \approx 0.05$, attention $h \approx 1$ (uniform), normalization $n \approx 0.32$ (GroupNorm washout). Their product $\kappa \approx 0.016$ matches the observed residual style coefficient. The ratio $w_{\text{SWD}}/w_{\text{FM}} = 0.01/1.0$ gives $\sqrt{w_{\text{SWD}}/w_{\text{FM}}} = 0.1 \gg \kappa$, so condition (i) holds. The attractor is multiplicative, so removing any single mechanism (e.g., increasing the gate initialization) is cancelled by the others. The Euclidean regime is therefore structurally bounded at CLIP-S ≈ 0.7175 . No reweighting of the three attenuations escapes the manifold, because their product stays below the style-gradient threshold.

Proposition 1 (Wavelet escape). *Under the spectral wavelet decomposition, the content gradient $\nabla_\theta \mathcal{L}_{\text{FM}}$ flows through the LL subband and the style gradient $\nabla_\theta \mathcal{L}_{\text{SWD}}$ flows through LH/HL/HH. By Proposition 2 these subspaces are orthogonal, so the Hessian $\nabla^2 \mathcal{L}$ block-diagonalizes: $\nabla^2 \mathcal{L} = \text{blkdiag}(\nabla_{\text{LL}}^2 \mathcal{L}_{\text{FM}}, \nabla_{\text{HF}}^2 \mathcal{L}_{\text{SWD}})$. The cross term $\nabla_{\text{LL, HF}}^2 \mathcal{L}$ vanishes, condition (ii) of Theorem 1 no longer binds the HF block, and $\mathcal{M}_0$ ceases to be a critical manifold of $\mathcal{L}_{\text{wavelet}}$.*

The wavelet decomposition of the next section provides this escape.

3.3 Spectral Decomposition

We apply a single-level 2D Haar wavelet transform to z_t , producing four subbands at half resolution. For each 2×2 patch (a, b, c, d) :

$$\begin{aligned} \text{LL} &= \frac{1}{2}(a + b + c + d), & \text{LH} &= \frac{1}{2}(a + b - c - d), \\ \text{HL} &= \frac{1}{2}(a - b + c - d), & \text{HH} &= \frac{1}{2}(a - b - c + d). \end{aligned} \quad (5)$$

The Haar matrix is orthogonal: $LL^2 + LH^2 + HL^2 + HH^2 = a^2 + b^2 + c^2 + d^2$. Reconstruction IDWT($\cdot$) is exact. A multi-level decomposition recurses on LL; we use three levels, the empirical sweet spot. Four levels loses positional information and degrades CLIP-S by 0.0003; two levels loses 0.0018.

Lemma 1 (Wavelet isometry and Parseval identity). *Let $\mathcal{W} : \mathbb{R}^{C \times H \times W} \rightarrow \mathbb{R}^{C \times H/2 \times W/2 \times \{4 \text{ subbands}\}}$ denote the single-level 2D Haar DWT. The Haar kernel matrix $H_2 = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}$ satisfies $H_2^\top H_2 = I_2$, so $\mathcal{W}$ is an isometry: $\|\mathcal{W}(z)\|_F^2 = \|z\|_F^2$. Consequently, for any partition of the four subbands into two disjoint sets A, B (e.g. $A = \{LL\}$, $B = \{LH, HL, HH\}$), $\|z\|_F^2 = \|A\|_F^2 + \|B\|_F^2$ with no cross term.*

Proposition 2 (Spectral content–style decoupling). *Let $z = z^{\text{cnt}} + z^{\text{sty}}$ be a decomposition in which the content signal concentrates its energy in LL, i.e. $\|z_{LL}^{\text{cnt}}\|_F^2 \geq (1 - \epsilon)\|z^{\text{cnt}}\|_F^2$ for some $\epsilon \in [0, 1]$. Then any transform T acting only on the high-frequency subbands $\{LH, HL, HH\}$ perturbs the content energy by at most $\epsilon \|z^{\text{cnt}}\|_F^2$:*

$$\left| \|T(z) - z\|_F^2 - \|T(z^{\text{sty}}) - z^{\text{sty}}\|_F^2 \right| \leq \epsilon \|z^{\text{cnt}}\|_F^2.$$

Proof. By Lemma 1 the LL and HF subspaces are orthogonal, so $\|T(z) - z\|_F^2 = \|T(z^{\text{sty}}) - z^{\text{sty}}\|_F^2 + \|T(z_{\text{HF}}^{\text{cnt}}) - z_{\text{HF}}^{\text{cnt}}\|_F^2$. The second term is bounded by $\|z_{\text{HF}}^{\text{cnt}}\|_F^2 \leq \epsilon \|z^{\text{cnt}}\|_F^2$ by the energy-concentration assumption. $\square$ $\square$

Subband semantics. LL carries global tone and object layout; LH/HL carry directional edges and brushwork; HH carries fine noise. In a controlled ablation, applying the LL velocity at inference contributes +0.0141 CLIP-S, while training the LL velocity contributes −0.0091 LPIPS via a gradient-backbone interaction. The two are decoupled, which motivates the design below.

3.4 Base Locking

We set the LL velocity loss weight $w_{\text{LL}} = 0$ during training but still apply the (small, freely drifting) predicted LL velocity at inference:

$$\mathcal{L}(\theta) = \sum_{c \in \{LH, HL, HH\}} w_c \mathbb{E}_t \|v_\theta^c(z_t, t, s) - u_t^c\|_2^2, \quad (6)$$

with $w_{\text{LH}} = w_{\text{HL}} = 1$ and $w_{\text{HH}} = 0$ (the HH head is removed entirely; in ablation, training vs. not training HH differs by $\Delta \text{CLIP-S} = \pm 0.0001$). This *base locking* choice is counter-intuitive: naively one would freeze LL at inference to “protect content.” We tried that (Table 3, “LL lock at inference”): it costs 0.0141 CLIP-S. LL also carries global style (color palette, lighting), and freezing it removes the only path by which the model can shift global tone. Letting LL drift freely during training but applying the drift at inference gives both the content anchor and the global-style channel. The content anchor holds because LL is barely trained, so v_θ^{LL} is small.

3.5 Fiber Flow: Per-Subband Style Injection

After the ODE reaches $\hat{z}_s$, we inject style by matching per-subband statistics. Write the endpoint latent and a reference style latent as

$$\hat{z}_s = \text{IDWT}(\hat{\ell}, \hat{\text{lh}}, \hat{\text{hl}}, \hat{\text{hh}}), \quad z_s^{\text{ref}} = \text{IDWT}(\ell^{\text{ref}}, \text{lh}^{\text{ref}}, \text{hl}^{\text{ref}}, \text{hh}^{\text{ref}}). \quad (7)$$

For each high-frequency subband $c \in \{\text{lh}, \text{hl}, \text{hh}\}$, we apply a whitening-and-coloring transform:

$$\text{WCT}(\hat{c}, c^{\text{ref}}) = \Sigma_{c^{\text{ref}}}^{1/2} \Sigma_{\hat{c}}^{-1/2} (\hat{c} - \mu_{\hat{c}}) + \mu_{c^{\text{ref}}}, \quad (8)$$

where μ, Σ are channel-wise mean and covariance, regularized as $\Sigma + 10^{-5}I$ for numerical stability. The LL subband is left unchanged: $\ell^{\text{new}} = \hat{\ell}$. The final latent is $\text{IDWT}(\ell^{\text{new}}, \text{lh}^{\text{new}}, \text{hl}^{\text{new}}, \text{hh}^{\text{new}})$.

This is a *fiber flow*: LL is the base manifold (content anchor), and the three high-frequency subbands are fibers above each LL point. Style transfer happens entirely in the fibers; the base is invariant. Because the Haar transform is orthogonal, per-fiber statistics matching is equivalent to per-frequency-band style injection in the original latent, with no cross-band leakage.

Definition 2 (Spectral fiber bundle). *The total space $E = \mathbb{R}^{C \times H \times W}$ is equipped with the projection $\pi : E \rightarrow B$ that returns the LL subband, $\pi(z) = z^{\text{LL}} \in B = \mathbb{R}^{C \times H/2 \times W/2}$. The fiber above a base point $b \in B$ is $F_b = \pi^{-1}(b) = \{z : z^{\text{LL}} = b\} \cong \mathbb{R}^{3C \times H/2 \times W/2}$, parameterized by the three high-frequency subbands (LH, HL, HH). A fiber-preserving map $T : E \rightarrow E$ satisfies $\pi(T(z)) = \pi(z)$ for all z .*

Proposition 3 (WCT is fiber-preserving). *The per-subband whitening-and-coloring transform WCT_c on each high-frequency subband $c \in \{\text{lh}, \text{hl}, \text{hh}\}$ leaves LL unchanged, hence $\pi(\text{WCT}(\hat{z}_s)) = \pi(\hat{z}_s) = \hat{\ell}$. The content anchor (LL) is an invariant of the style injection; style transfer is a vertical map along the fibers F_b .*

3.6 End-of-Trajectory Injection

A subtle failure of per-step style injection is that the residual content fraction after n AdaIN steps with blend α is $r_n = (1 - \alpha)^n$. For $n = 8$ and $\alpha = 0.5$, $r_n \approx 0.0039$, so $\alpha = 0.5$ and $\alpha = 1.0$ are indistinguishable. We therefore inject style *only at the last ODE step*: the first $n - 1$ steps are pure ODE integration, and WCT is applied once at $t = 1$. This restores α as a meaningful trade-off knob (verified by a monotone α -ablation). It also matches the Schrödinger bridge view, in which style is a terminal constraint, not a per-step perturbation.

Proposition 4 (α -identifiability under endpoint injection). *Let $r_n(\alpha) = (1 - \alpha)^n$ be the content residual after n per-step injections with blend $\alpha \in (0, 1]$. Two blends $\alpha_1 < \alpha_2$ are δ -indistinguishable if $|r_n(\alpha_1) - r_n(\alpha_2)| < \delta$. Under per-step injection with $n \geq 2$, the set of δ -indistinguishable pairs covers $[0, 1]^2$ for any $\delta > (1 - \alpha_{\min})^n$. Under endpoint injection ($n = 1$), $r_1(\alpha) = 1 - \alpha$ is injective in α , so all pairs are distinguishable for any $\delta < 1$.*

Proof. For $n \geq 2$, r_n is convex and decreasing on $[0, 1]$ with $r_n(0) = 1$, $r_n(1) = 0$; its Lipschitz constant is $n(1 -$

$\alpha_{\min})^{n-1}$, which exceeds 1 for small $\alpha_{\min}$, collapsing the resolution. For $n = 1$, $r_1(\alpha) = 1 - \alpha$ is affine with slope -1 , hence injective. $\square$ $\square$

3.7 Stochastic Frequency Routing

The cross-attention layer that routes style tokens into the velocity field can operate on the full latent (Euclidean query) or on the high-frequency tokens only (wavelet query). At inference we always use the wavelet query. At training time, we draw a Bernoulli variable $B \sim \text{Bernoulli}(p)$ per step and use the wavelet query with probability p , the full query otherwise. We set $p = 0.8$.

The rationale is a train-test distribution mismatch: if $p = 1.0$ throughout training, the style encoder over-specializes to high-frequency statistics and loses global palette information. If $p = 0.0$ (wavelet only at inference), the query projection has never seen wavelet coefficients and the route fails (CLIP-S drops to 0.7061). The sweep $p \in \{0.0, 0.5, 0.8, 1.0\}$ gives $\text{CLIP-S} \in \{0.7061, 0.7083, 0.7213, 0.7226\}$ and $\text{LPIPS} \in \{0.2606, 0.2480, 0.2868, 0.3068\}$. $p = 0.8$ is the unique point where both metrics pass the dual threshold.

3.8 Architecture and Training

The backbone is four residual blocks at $C=64$ with AdaIN conditioning. It processes the concatenated subbands $[z^{\text{LL}}, z^{\text{LH}}, z^{\text{HL}}] \in \mathbb{R}^{12 \times 16 \times 16}$ (HH is dropped, see Section 3.4). Per-subband 3×3 convolutional heads output $v^{\text{LL}}, v^{\text{LH}}, v^{\text{HL}} \in \mathbb{R}^{4 \times 16 \times 16}$. The style encoder is a frozen DINOv2 backbone with a learnable per-domain token bank (5 tokens $\times$ 64 dims). Total trainable parameters: 903,248. We train for 5 epochs with batch size 24, Adam at $\text{lr } 10^{-4}$, on a single RTX 3060 (12 GB). End-to-end training time is 3 minutes 5 seconds including data loading and checkpoint saving.

4 Experiments

4.1 Setup

Dataset. We use the Distinct5-WikiArt-512 benchmark: five artistic domains — Early Renaissance, Impressionism, Minimalism, Rococo, Ukiyo-e. Each domain has 3,600 training images and 150 test images. Inputs are SDXL VAE latents at $4 \times 32 \times 32$ (512×512 images). Training uses 18,000 latents; evaluation uses 750 pairs (5 source domains $\times$ 5 target domains $\times$ 30 source images).

Metrics. CLIP-S is the cosine similarity between CLIP ViT-B/32 embeddings of the output and a target-style prototype (higher is better). LPIPS is the AlexNet perceptual distance between output and content image (lower is better). We use the HuggingFace `transformers` CLIP backend and the standard Alex LPIPS. All baselines are evaluated under this unified protocol.

Baselines. Twelve methods: Identity (no transfer), AdaIN, WCT (VGG19), SD-Turbo, SDEdit at $s \in \{0.35, 0.40\}$, StyleID, CUT, SaMST, SaMam, Seedream 4.5 (commercial API), and our method. All trainable baselines (CUT, SaMST, SaMam) are trained to convergence on the same Distinct5

Table 1: Main comparison on the 750-pair Distinct5-WikiArt-512 protocol (5 source domains $\times$ 5 target domains $\times$ 30 source images, including 150 identity pairs where source = target domain). All trainable baselines trained to convergence on the same 18K-image set. $\uparrow$: higher better, $\downarrow$: lower better. Best non-Identity result in **bold**. Bootstrap analysis and per-pair caliber details in the supplement.

#	Method	Category	CLIP-S $\uparrow$	LPIPS $\downarrow$	Train (min)
1	Identity	—	0.6933	0.0000	0
2	AdaIN	classical-inf	0.6679	0.7425	0
3	WCT (VGG19)	classical-inf	0.7063	0.6348	0
4	SD-Turbo	diffusion-inf	0.6933	0.0033	0
5	SDEdit $s=0.35$	diffusion-sweep	0.7797	0.4508	0
6	SDEdit $s=0.40$	diffusion-sweep	0.7934	0.4826	0
7	StyleID	diffusion-inf	<u>0.8223</u>	0.5523	0
8	CUT	gan-train	0.7137	0.3743	322.6
9	SaMST	mamba-train	0.6183	0.7490	39.5
10	SaMam	mamba-train	0.5816	<u>0.2434</u>	436
11	Seedream 4.5	commercial-api	0.7198	0.4767	—
12	SFM (Ours)	spectral-flow	0.7213	0.2868	3.08

training set. SaMam is trained for 20,000 steps (436 minutes); CUT for 322.6 minutes.

4.2 Main Results

Table 1 reports the main comparison. Three observations stand out.

SFM dual-beats SaMam. On the 750-pair protocol (which includes 150 identity pairs), SaMam’s CLIP-S (0.5816) is below the Identity baseline (0.6933). This means SaMam performs content reconstruction rather than style transfer. Its low LPIPS (0.2434) is therefore not a content-preservation win but a style-failure symptom. SFM exceeds SaMam on CLIP-S by $+0.1397$ while training $141 \times$ faster (3.08 min vs. 436 min). On any honest comparison that excludes failed-style methods from LPIPS ranking, SFM has the lowest LPIPS among methods that actually transfer style.

We report Bootstrap 95% confidence intervals for the style-transfer methods with available per-pair predictions (SaMST, Seedream 4.5) in the supplement. SFM and SaMam per-pair predictions will be released with the code.

SFM dual-beats Seedream 4.5. The commercial API achieves CLIP-S 0.7198 with LPIPS 0.4767. SFM matches it on CLIP-S ($+0.0015$) and beats it on LPIPS by -0.1899 . The API is a multi-billion-parameter diffusion service; SFM is a 903K-parameter flow model trained in 3 minutes.

SFM beats Identity on style, matches it on content. Identity (no transfer) is the content-preservation ceiling (LPIPS=0). SFM’s LPIPS of 0.2868 is the style-transfer cost; its CLIP-S of 0.7213 exceeds Identity by $+0.0280$, confirming that real style is transferred. By contrast, AdaIN/WCT/SaMST all have CLIP-S at or below Identity, meaning they fail to transfer style at all on this benchmark.

Why high-CLIP-S baselines are not winners. StyleID (0.8223) and SDEdit (0.7934) achieve the highest CLIP-S




Figure 2: CLIP-S vs. 1-LPIPS across all 12 baselines. SFM is the upper-left point. Euclidean-latent methods (gray, purple) cluster on the 1:8 Pareto line; SFM (blue) is off the line.



Figure 3: Qualitative comparison. SFM transfers brushwork and texture (Real Style) while preserving content. StyleID tears content; AdaIN applies only a color filter; SaMam re-constructs content with no style.

but at $\text{LPIPS} > 0.45$, indicating severe content tearing. The Pareto frontier in Figure 2 shows that these points are dominated. SFM sits in the upper-left (high CLIP-S, low LPIPS) region that no Euclidean method reaches.

4.3 Efficiency: 3 Minutes on an RTX 3060

SFM’s training cost is dominated by 5 epochs over 18,000 latents at batch 24, which is 3,750 forward-backward passes. On an RTX 3060, this takes 1 min 12 sec of pure training time. End-to-end it takes 3 min 5 sec, including VAE cache loading, optimizer initialization, and checkpoint save. The reason the loss landscape is benign is structural: with LL locked, the velocity field only needs to learn brushwork, which is a low-complexity target. Euclidean-latent baselines spend most of their compute reconciling content reconstruction with style injection, which is a high-dimensional non-convex problem.

Convergence analysis. The flow-matching loss on the high-frequency subbands is approximately convex. To see why, write the loss as $\mathcal{L}_{\text{HF}}(\theta) = \mathbb{E}_t \|v_\theta^{\text{HF}}(z_t, t, s) - u_t^{\text{HF}}\|^2$. The target $u_t^{\text{HF}} = z_1^{\text{HF}} - z_0^{\text{HF}}$ is a fixed function of the data. By Proposition 2, the HF subspace carries $\leq \epsilon$ of the content energy, so u_t^{HF} is dominated by style. The velocity field v_θ^{HF} is a residual network with Lipschitz-continuous activations,

Table 2: Training efficiency. SFM trains in 3 minutes on a consumer-grade RTX 3060 (12 GB), $141\times$ faster than SaMam and $105\times$ faster than CUT.

Method	Params	Train time	Hardware	Speedup
CUT	28.4M	322.6 min	RTX 3090	$105\times$
SaMST	6.6M	39.5 min	RTX 3090	$12.8\times$
SaMam	5.1M	436 min	RTX 3090	$141\times$
SFM (Ours)	0.903M	3.08 min	RTX 3060	—

so $\mathcal{L}_{\text{HF}}$ is locally strongly convex near the optimum. By contrast, the Euclidean loss $\mathcal{L}_{\text{FM}} = \mathbb{E}_t \|v_\theta(z_t, t, s) - u_t\|^2$ has a target $u_t = z_1 - z_0$ that mixes content and style. The two gradient components $\nabla_z v_\theta \cdot \nabla_{\text{content}}$ and $\nabla_s v_\theta \cdot \nabla_{\text{style}}$ point in conflicting directions on the degeneration manifold (Theorem 1), producing a saddle-shaped landscape. Empirically, the Euclidean baseline requires 50–100 epochs to reach a stationary point still below the identity baseline. SFM reaches its optimum in 3 epochs and is stable by epoch 5.

Latent vs. pixel space. A direct pixel-space comparison is infeasible on a 12 GB consumer GPU. Pixel-space flow matching on 512×512 images requires roughly $192\times$ the compute of latent-space training. Attention complexity scales quadratically with sequence length (a $65536\times$ increase for full-resolution attention). Prior work (Rombach *et al.* 2022) demonstrated that latent diffusion achieves superior sample quality to pixel-space diffusion at equal compute. The semantic abstraction of VAE latents is particularly beneficial for style transfer, where low-level pixel statistics entangle with content. Our use of latent space is therefore both a computational necessity and a representational advantage.

4.4 Ablation: Component Analysis

Table 3 reports component ablations that remove or replace one component at a time from the SFM baseline.

The 1:8 trade-off. The capacity knobs (depth, dim, gate) all move the operating point along a single line with slope ≈ 8 . Reducing capacity improves LPIPS by $8\times$ the CLIP-S cost. The same line is traced by α , schedule shape, and training duration. This is the Pareto deadlock of Section 3.2: none of these knobs breaks the frontier. The only structural degrees of freedom that do are (i) the wavelet decomposition itself, (ii) the solver order (Euler→Heun), and (iii) the stochastic routing probability p .

Why capacity maps to the 1:8 line. On the degeneration manifold $\mathcal{M}_0$ (Definition 1), the velocity field factorizes as $v_\theta(z, t, s) = \bar{v}(z, t) + \kappa(\theta) \delta v_s(z, t)$. Here $\kappa(\theta) = g h n$ is the multiplicative attenuation of Theorem 1. The content-preservation term $\mathcal{L}_{\text{FM}}$ is independent of κ ; the style term $\mathcal{L}_{\text{SWD}}$ scales as κ^2 . Reducing capacity (e.g., depth $4 \rightarrow 2$) reduces $\|\delta v_s\|$ by a factor β . This reduces CLIP-S by $O(\beta)$ and reduces LPIPS by $O(\beta)$ as well (less style injection $\Rightarrow$ less content distortion). The ratio $\Delta \text{LPIPS} / \Delta \text{CLIP-S} \approx 8$ is therefore the ratio of the two loss curvatures on $\mathcal{M}_0$, not

Table 3: Component ablation. Each row removes or replaces one component of SFM. Δ CLIP-S and Δ LPIPS are relative to the SFM baseline (0.7213 / 0.2868).

Variant	CLIP-S	LPIPS	Δ CLIP-S	Note
SFM (full)	0.7213	0.2868	—	baseline
<i>Wavelet decomposition (the active ingredient)</i>				
No DWT (Euclidean)	0.7117	0.2994	−0.0096	attractor active
DWT level 1 (vs. 3)	0.7261	0.3296	+0.0048	under-decomposition
DWT level 4 (vs. 3)	0.7316	0.3461	+0.0103	over-decomposition
Daubechies-2 (vs. Haar)	0.7258	0.3288	+0.0045	basis function
<i>Base locking</i>				
LL lock at inference	0.7174	0.3372	−0.0039	kills global style channel
$w_{LL} = 1.0$ at train	0.7174	0.2774	−0.0039	content-heavy color and HH head rendering
Train HH head	0.7213	0.2868	±0.0001	HH head rendering
<i>Style injection</i>				
Per-step AdaIN (not EOTA)	0.7361	0.3843	+0.0148	α -decay, LPIPS explodes
AdaIN (diagonal) vs. WCT	0.7266	0.3402	+0.0053	loses channel correlation
Inject LL too	0.7215	0.3856	+0.0002	color shift propagates
<i>Stochastic routing</i>				
$p = 0.0$ (Euclidean train)	0.7061	0.2606	−0.0152	train/test mismatch
$p = 0.5$	0.7083	0.2480	−0.0130	under-trained model
$p = 1.0$ (always DWT)	0.7226	0.3068	+0.0013	style encoder specialized
<i>Solver</i>				
Euler (vs. Heun)	0.7210	0.3129	−0.0003	structural BFGS loss
RK4 (vs. Heun)	0.7265	0.3235	+0.0052	saturated, no gain
<i>Capacity (Pareto-mapping knobs)</i>				
depth=2 (vs. 4)	0.7172	0.2705	−0.0041	1:8 trade-off
dim=32 (vs. 64)	0.7153	0.2705	−0.0060	1:8 trade-off
gate=0 (no gate)	0.7080	0.2641	−0.0133	1:8 trade-off
depth=6 (vs. 4)	NaN	NaN	—	WCT eigh unstable

a property of any specific knob. This is why *every* capacity knob traces the same line: they all perturb κ , and κ is the only degree of freedom on $\mathcal{M}_0$. The wavelet decomposition escapes this by moving the model off $\mathcal{M}_0$ entirely (Proposition 1), making κ irrelevant.

5 Discussion and Limitations

Why does it train in 3 minutes? Not because the model is small (903K params), and not because the data is small (18K latents). It trains in 3 minutes because the loss landscape is convex once LL is locked. In the Euclidean regime, the model must reconcile two contradictory gradients: content preservation pulls toward z_0 ; style injection pulls toward z_1 . The reconciliation is a non-convex negotiation that takes hours. In the wavelet regime, the LL gradient is zero by construction, and the LH/HL/HH gradients are aligned (they all point toward style). The model has nothing to negotiate; it just fits brushwork.

Why is it “real” style transfer? AdaIN and WCT in Euclidean space shift the global mean and covariance of the entire feature map, which manifests as a color filter. SFM applies WCT per-subband on LH/HL/HH only, leaving LL (global tone and layout) intact. The visible effect is that brushstroke direction, edge sharpness, and texture granular-

ity are transferred while the content image’s color palette is preserved. This is what we mean by “Real Style”: not a color filter, but a texture-level transfer.

Limitations. (1) The style is domain-level, not per-image: SFM transfers the *style distribution* of a domain, not the unique brushwork of a single painting. Per-image style would require a style encoder conditioned on a single reference, which we leave to future work. (2) The Haar wavelet is the simplest orthogonal transform; biorthogonal or learned wavelets might give finer control. (3) When the content image has strong high-frequency structure (text, fine patterns), the high-frequency style injection may overwrite it. (4) Five domains is a limited evaluation; extension to sketch, watercolor, and 3D-rendered styles is left to future work. (5) Transferring to Minimalism, whose style is largely low-frequency color blocking, is harder because base locking limits global tone shift. This is the intended trade-off for real texture transfer.

6 Conclusion

We have shown that the Pareto deadlock in unsupervised domain style transfer is a structural consequence of content specialization in Euclidean space. Spectral Flow Matching dissolves it by moving transport to the orthogonal wavelet domain: 3 minutes on a consumer RTX 3060, 903K parameters, dual-beating a 436-minute Mamba baseline and a commercial diffusion API. The community does not need larger models or more paired data; it needs the right coordinate system.

Ethics Statement

Style transfer tools can be misused to produce deepfake art or to misattribute stylistic authorship. Our method operates at the domain level (e.g., “Impressionism”) rather than the per-artist level, which reduces the risk of impersonating a specific living artist. We do not condition on individual reference paintings. The dataset is sourced from WikiArt, which consists of public-domain and historically significant artworks; no personal data is used. We encourage downstream users to disclose when an image has been style-transferred and to respect the licensing terms of the source content.

Reproducibility Statement

All experiments use public data (WikiArt) and a public VAE (SDXL VAE). The training configuration, hyperparameters, and evaluation protocol are described in Section 4 and the supplement. Training takes 3 minutes on a single RTX 3060 (12 GB); evaluation takes 97 seconds on the same hardware. The 903K-parameter checkpoint is 3.5 MB. We will release the code, the checkpoint, and the 750-pair evaluation protocol upon publication.

AAAI Reproducibility Checklist

1. This paper claims the following contributions: (1) diagnosis of the Pareto deadlock and degeneration attractor (Section 3.2); (2) Spectral Flow Matching method (Section 3); (3) experimental validation on 750 pairs (Section 4).

2. The dataset (Distinct5-WikiArt-512) is publicly available; the evaluation protocol (CLIP-S, LPIPS, 750 pairs) is described in Section 4.
3. Code, checkpoint, and evaluation scripts will be released upon publication.
4. Training requires a single 12 GB GPU (RTX 3060); 5 epochs; 3 minutes end-to-end.
5. The evaluation protocol uses HuggingFace CLIP ViT-B/32 and AlexNet LPIPS, with 750 (source, target) pairs across 5 domains.
6. All baseline numbers (Table 1) are evaluated under the same protocol; SaMam and CUT are trained to convergence on the same Distinct5 training set.
7. Statistical significance: single-seed results are reported in Table 1. Bootstrap 95% confidence intervals for SaMST and Seedream 4.5 (methods with available per-pair predictions) are in the supplement. Multi-seed variance for SFM is documented in the supplement.

References

- Jiwoo Chung, Sangeek Hyun, and Jae-Pil Heo. Style injection in diffusion: A training-free approach for adapting large-scale diffusion models for style transfer. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024.
- Ingrid Daubechies. Ten lectures on wavelets. *SIAM*, 1992.
- Yingying Deng, Fan Tang, Weiming Dong, Wen Sun, Feiyue Huang, and Changsheng Xu. Stytr2: Image style transfer with transformers. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2022.
- Leon A. Gatys, Alexander S. Ecker, and Matthias Bethge. A neural algorithm of artistic style. In *arXiv preprint arXiv:1508.06576*, 2015.
- Leon A. Gatys, Alexander S. Ecker, and Matthias Bethge. Image style transfer using convolutional neural networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016.
- Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv preprint arXiv:2312.00752*, 2023.
- Judy Hoffman, Eric Tzeng, Taesung Park, Jun-Yan Zhu, Kate Kate, Kate Saenko, Trevor Darrell, and Alexei A. Efros. Cycada: Cycle-consistent adversarial domain adaptation. In *International Conference on Machine Learning*, 2018.
- Xun Huang and Serge Belongie. Arbitrary style transfer in real-time with adaptive instance normalization. In *Proceedings of the IEEE International Conference on Computer Vision*, 2017.
- Yujun Li, Chen Chen, Jimei Yang, Lijuan Wang, Zhe Lin, and Yi Yang. Universal style transfer via feature transforms. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2017.
- Yaron Lipman, Ricky T. Q. Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow matching for generative modeling. *arXiv preprint arXiv:2210.02747*, 2022.
- Yahui Liu et al. Pluggable style representation learning for multi-style transfer. In *Proceedings of the Asian Conference on Computer Vision*, 2024.
- Yahui Liu et al. Samam: Style-aware mamba for style transfer. *arXiv preprint arXiv:2403.11887*, 2024.
- Chenlin Meng, Yutong He, Yang Song, Jiaming Song, James Blackwell, Jiazhao Jia, and Stefano Ermon. Sdedit: Guided image synthesis and editing with stochastic differential equations. In *Proceedings of the IEEE International Conference on Computer Vision*, 2021.
- Dae Young Park and Kwang Hee Lee. Arbitrary style transfer with style-attentional networks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019.
- Taesung Park, Alexei A. Efros, Richard Zhang, and Jun-Yan Zhu. Contrastive learning for unpaired image-to-image translation. In *Proceedings of the European Conference on Computer Vision*, 2020.
- Thu Phung, Prasanna Ghosh, and Long Tran. Wavelet diffusion models are fast and scalable image generators. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023.
- Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2022.