

Affordable Real Style Transfer: Training Wavelet-Decoupled Rectified Flow on a Single RTX 3060 in Minutes

Anonymous submission

Abstract

Art-to-art style transfer is easy to overclaim because the source image is already an artwork. If the unchanged source already resembles the target domain, high style similarity can come from doing almost nothing. We formalize this issue with identical-image transfer (IDT). The unchanged source defines a no-op floor. A method below that floor has not moved in the requested direction.

We then analyze why compact flow models often fall into this regime. In Euclidean latent coordinates, low-frequency layout motion and high-frequency texture motion share one basis. The flow objective can then favor structure preservation and suppress style-conditioned transport. We state a sufficient local condition for this style-suppression regime and turn it into ablation-level predictions.

We propose WD-VF, a wavelet-decoupled rectified flow in the latent space of the Stable Diffusion v1.5 EMA VAE. A single-level Haar transform separates LL, LH, HL, and HH subbands.

WD-VF weakens LL supervision, routes style queries through high-frequency bands, and applies whitening and coloring only at the endpoint. This keeps global structure stable while concentrating model capacity on stroke and texture transfer.

On Distinct5-WikiArt, the 903K-parameter WD-VF model reaches CLIP-S=0.7213 and LPIPS=0.2868. It clears the IDT floor of 0.6933 and exceeds SaMam by +0.1397 CLIP-S. Relative to Seedream 4.5, it reduces LPIPS by 39.8% while staying in the same ArtFID range. Training takes 3.08 minutes on a single RTX 3060 (12 GB). Generation plus VAE decode takes 83.7 seconds for 750 outputs, or 111.7 ms/image. Furthermore, this decoupled geometry enables lightweight few-shot adaptation. Adding new styles requires updating only a compact style-memory module without retraining the core transport flow, preserving high-quality transfer. The ablations follow the predicted pattern: removing DWT weakens the frontier, route mismatch lowers transfer quality, and repeated style injection raises LPIPS.




Figure 1: IDT-calibrated comparison on Distinct5-WikiArt. Left: CLIP-S versus $1 - \text{LPIPS}$, with bubble area encoding training time. The red points mark two WD-VF operating points. The dashed line marks the IDT floor (0.6933), which SaMam falls below. Right: target-pooled ArtFID on the same benchmark, with numbers inside the bars indicating training or response time.

1 Introduction

Art-to-art style transfer has an unusually forgiving failure mode. The source image is already an artwork, so doing almost nothing can still produce a high style score. If the unchanged source already lies close to the target domain,

raw CLIP similarity does not prove that the method moved in the requested direction. We therefore calibrate this setting with identical-image transfer (IDT). The unchanged source defines a no-op floor. Real transfer must clear that floor.

This calibration changes the ranking on Distinct5-WikiArt. We build the benchmark from the five WikiArt styles with

the lowest IDT CLIP-S in a broader style pool. This makes accidental no-op success harder. Under the unified 750-pair protocol, SaMam reaches LPIPS 0.2434 yet only CLIP-S 0.5816. That score is far below the IDT floor of 0.6933. The result is not strong transfer with low damage. It is near-reconstruction in the wrong direction.

We trace this failure to the geometry of Euclidean latent flow. In standard latent coordinates, the rectified-flow target $u_t = z_1 - z_0$ mixes low-frequency layout motion with high-frequency texture motion. The flow loss can then reward structure preservation more strongly than style-conditioned motion. Optimization is pulled toward a style-suppression regime. Theorem 1 formalizes one sufficient local condition for that regime. The theorem also yields concrete predictions. Removing the coordinate change should weaken the CLIP-S/LPIPS frontier. Route mismatch should lower transfer quality. Repeated style injection should increase content damage.

WD-VF solves the problem by changing coordinates rather than enlarging the generator. It applies a single-level Haar transform, weakens LL supervision, routes style queries through the high-frequency bands, and stylizes only the endpoint subbands. In the ideal routed limit, the active supervised residual reduces to two orthogonal high-frequency bands. This removes the dominant low-frequency conflict from both the style-query path and most of the supervised transport.

The resulting model is explicit and small. It uses 903K parameters. It reaches CLIP-S 0.7213 and LPIPS 0.2868, which clears the IDT floor. Among trained local baselines in this benchmark, it gives the strongest positive-IDT result. Training takes 3.08 minutes on a single RTX 3060. Generation plus VAE decode takes 83.7 seconds for 750 outputs, or 111.7 ms/image.

Contributions. (1) We introduce IDT-calibrated evaluation for art-to-art transfer. It distinguishes real target-direction transfer from no-op behavior on artwork inputs. (2) We analyze a local style-suppression regime in Euclidean latent flow. Theorem 1 states a sufficient curvature condition and yields ablation-level predictions. (3) We propose **WD-VF**, a wavelet-decoupled rectified flow with weak LL supervision, high-frequency query routing, and endpoint subband alignment. In the ideal routed limit, the active supervised transport reduces to two orthogonal high-frequency bands (Proposition 2). (4) We show that this geometry gives an efficient real-transfer regime. WD-VF uses 903K parameters and trains in 3.08 minutes. Crucially, the decoupled transport is highly extensible: we demonstrate that the model can generalize to unseen styles via lightweight style-memory updating, eliminating the need for full retraining.

2 Related Work

Exemplar methods. Classical exemplar transfer uses a target image at inference time (Gatys, Ecker, and Bethge 2015, 2016; Huang and Belongie 2017; Li et al. 2017; Deng et al. 2022; Huang et al. 2024). These methods match target statistics or target correspondences. Our setting uses only the source image and a target style identity. The task is therefore

domain-conditional translation rather than exemplar stylization.

Domain-conditional methods. The closest learned baselines are compact multi-style systems such as SaMST and SaMam (Liu et al. 2024a,b). Larger diffusion-based editors such as StyleID and SDEdit, together with commercial systems such as Seedream, define the high-compute end of the same benchmark range (Meng et al. 2021; Chung, Hyun, and Heo 2024; Zhang et al. 2024). WD-VF targets the same interface at much lower cost.

Wavelets and evaluation. Prior wavelet methods mainly use multi-scale decompositions as feature-processing modules (Daubechies 1992; Phung, Ghosh, and Tran 2023). We instead use a wavelet basis as a coordinate change that separates low-frequency structure from high-frequency style motion. We pair this design with IDT calibration because the source image is already an artwork. The main metrics are CLIP-S, LPIPS (Zhang et al. 2018), and the gap to the IDT floor. ArtFID (Wright and Ommer 2022) serves as an auxiliary check.

3 Theory and Method

3.1 Setup

Let $x_c \in \mathbb{R}^{H \times W \times 3}$ be a content image and $s \in \{1, \dots, K\}$ a target style domain. We work in the latent space of the Stable Diffusion v1.5 EMA VAE. We write

$$z_0 = \mathcal{E}(x_c) \in \mathbb{R}^{4 \times 32 \times 32}. \quad (1)$$

During training, $z_1 \sim p_s$ is a latent sampled from the target domain. Rectified flow uses the straight segment

$$\begin{aligned} z_t &= (1-t)z_0 + tz_1, \\ u_t &= z_1 - z_0, \quad t \sim \mathcal{U}([0, 1]). \end{aligned} \quad (2)$$

At inference we use an 8-step Heun solver,

$$\begin{aligned} v_t &= v_\theta(z_t, t, s), \\ v'_t &= v_\theta(z_t + \Delta t v_t, t + \Delta t, s), \\ z_{t+\Delta t} &= z_t + \frac{\Delta t}{2} (v_t + v'_t), \end{aligned} \quad (3)$$

We then decode $\hat{x}_s = \mathcal{D}(\hat{z}_1)$. In our controls, Heun gives the best cost-quality tradeoff. Euler loses style. RK4 adds cost without improving the final result.

3.2 A Local Collapse Mechanism in Euclidean Latents

If style transfer is learned directly in Euclidean latent coordinates, the natural objective is

$$\mathcal{L}_{\text{euc}} = w_f \mathcal{L}_{\text{FM}} + w_s \mathcal{L}_{\text{sty}}, \quad (4)$$

where $\mathcal{L}_{\text{FM}} = \mathbb{E}_t \|\tilde{v}_\theta(z_t, t, s) - u_t\|_2^2$ and $\mathcal{L}_{\text{sty}}$ is any endpoint style-matching term. The difficulty is structural. The straight target $u_t = z_1 - z_0$ mixes low-frequency layout motion with high-frequency texture motion in the same coordinates.

Definition 1. Define

$$\mathcal{M}_0 = \{\theta : \tilde{v}_\theta(z, t, s) = \bar{v}_\theta(z, t) \text{ for all } z, t, s\}, \quad (5)$$

with $\bar{v}_\theta(z, t) = \mathbb{E}_s[\tilde{v}_\theta(z, t, s)]$.




Figure 2: WD-VF architecture. The latent is decomposed into LL, LH, HL, and HH subbands. The diagram highlights style query routing (LL is excluded). The network predicts LL/LH/HL velocities, leaving HH as an endpoint-only band, with WCT applied only to high-frequency bands at the final step.

The set $\mathcal{M}_0$ contains parameters whose field ignores the style label. This matches the low-LPIPS, sub-IDT regime observed empirically. The Euclidean baseline is well described by

$$\tilde{v}_\theta(z, t, s) = \bar{v}_\theta(z, t) + \kappa(\theta) \Delta v_\theta(z, t, s), \quad (6)$$

$$\mathbb{E}_s[\Delta v_\theta(z, t, s)] = 0,$$

where $\kappa(\theta)$ is the retained style amplitude. Points on $\mathcal{M}_0$ satisfy $\kappa(\theta) = 0$.

Theorem 1. *Let $\theta_0 \in \mathcal{M}_0$ be stationary in the normal directions:*

$$\nabla_\perp \mathcal{L}_{\text{FM}}(\theta_0) = 0, \quad \nabla_\perp \mathcal{L}_{\text{sty}}(\theta_0) = 0. \quad (7)$$

Suppose the normal curvatures satisfy

$$\nabla_\perp^2 \mathcal{L}_{\text{FM}}(\theta_0) \succeq \mu I, \quad \|\nabla_\perp^2 \mathcal{L}_{\text{sty}}(\theta_0)\|_{\text{op}} \leq \nu, \quad (8)$$

If $w_f \mu > w_s \nu$, then θ_0 is a strict local minimum of $\mathcal{L}_{\text{euc}}$ on the normal space $T_{\theta_0}^\perp \mathcal{M}_0$.

Proof. For any normal perturbation $\delta \in T_{\theta_0}^\perp \mathcal{M}_0$, the stationarity conditions (7) remove all first-order terms. A second-order expansion gives

$$\Delta \mathcal{L}_{\text{euc}}(\delta) = \frac{1}{2} \delta^\top \left(w_f \nabla_\perp^2 \mathcal{L}_{\text{FM}}(\theta_0) + w_s \nabla_\perp^2 \mathcal{L}_{\text{sty}}(\theta_0) \right) \delta + o(\|\delta\|^2). \quad (9)$$

Using (8),

$$\Delta \mathcal{L}_{\text{euc}}(\delta) \geq \frac{1}{2} (w_f \mu - w_s \nu) \|\delta\|^2 + o(\|\delta\|^2). \quad (10)$$

The coefficient of $\|\delta\|^2$ is positive. Therefore $\Delta \mathcal{L}_{\text{euc}}(\delta) > 0$ for all sufficiently small nonzero δ .

The two curvature bounds have a simple meaning. The flow term is locally steep in the normal directions. The style

term is locally flatter. When $w_f \mu > w_s \nu$, optimization is pulled toward $\mathcal{M}_0$ instead of a style-conditioned direction.

Theorem 1 is the first half of the paper’s causal chain. It identifies a concrete Euclidean failure regime and predicts what should happen when that regime is broken. If we remove the shared Euclidean coordinates, the CLIP-S/LPIPS frontier should improve. If a Euclidean fit remains trapped near $\mathcal{M}_0$, its retained style amplitude should stay small. In our Euclidean fit, the amplitude estimate is $\kappa \approx 0.016$. In Table 2, removing the wavelet coordinate change weakens the frontier exactly as predicted.

The theorem is local rather than global. It does not claim that every Euclidean flow must collapse. It states that near $\mathcal{M}_0$, flow matching can dominate endpoint style supervision in the normal directions. That is precisely the regime relevant to our benchmark: the Euclidean baseline has low LPIPS, sub-IDT CLIP-S, and a fitted style amplitude close to zero.

3.3 Wavelet Coordinates

WD-VF first changes coordinates. For a 2×2 patch (a, b, c, d) , the single-level Haar transform is

$$\mathcal{W}(z) = (\ell, h_1, h_2, h_3), \quad (11)$$

with

$$\begin{aligned} \ell &= \frac{1}{2}(a + b + c + d), & h_1 &= \frac{1}{2}(a + b - c - d), \\ h_2 &= \frac{1}{2}(a - b + c - d), & h_3 &= \frac{1}{2}(a - b - c + d). \end{aligned} \quad (12)$$

Here ℓ is the low-frequency coordinate and $h = (h_1, h_2, h_3)$ collects the high-frequency coordinates.

Lemma 1. *The transform is orthogonal, so*

$$\|z\|_F^2 = \|\ell\|_F^2 + \|h_1\|_F^2 + \|h_2\|_F^2 + \|h_3\|_F^2. \quad (13)$$

Definition 2. *We use ℓ as the low-frequency coordinate and (h_1, h_2, h_3) as the high-frequency coordinates. Equivalently, the projection map is $\pi(z) = \ell$.*

Proposition 1. Let $z = z_{\text{cnt}} + z_{\text{sty}}$ and write $\mathcal{W}(z_{\text{cnt}}) = (\ell_{\text{cnt}}, h_{\text{cnt}})$. If

$$\|h_{\text{cnt}}\|_F^2 \leq \epsilon \|z_{\text{cnt}}\|_F^2 \quad \text{for some } \epsilon \in [0, 1], \quad (14)$$

then any high-frequency-only map $T_\ell(\ell, h) = (\ell, \phi(h))$ satisfies

$$| \|T_\ell(z) - z\|_F^2 - \|\phi(h) - h\|_F^2 | \leq \epsilon \|z_{\text{cnt}}\|_F^2. \quad (15)$$

Proof. Lemma 1 removes all cross terms between ℓ and h . The only content perturbation left is the small content residue already present in h .

Proposition 1 states the geometric premise of WD-VF. When most content energy lies in LL, high-frequency transport changes texture much more than layout. This motivates weak LL supervision and high-frequency routing.

This premise is approximate rather than absolute. High-frequency bands can still contain some content residue, and some styles also change coarse palette. WD-VF therefore keeps a weak LL term instead of setting it to zero in the final model. The wavelet basis is useful because it makes the dominant conflict explicit and controllable.

3.4 Training Objective

Let

$$\begin{aligned} \mathcal{W}(z_t) &= (\ell_t, h_{1,t}, h_{2,t}, h_{3,t}), \\ \mathcal{W}(u_t) &= (u_\ell, u_1, u_2, u_3). \end{aligned} \quad (16)$$

The final WD-VF model receives $(\ell_t, h_{1,t}, h_{2,t}, h_{3,t})$ and predicts three velocity heads (v_ℓ, v_1, v_2) . The h_3 band stays in the state and is stylized at the endpoint. It has no learned velocity head.

The implemented training objective is

$$\mathcal{L}_{\text{impl}} = \mathbb{E}_t [\lambda_\ell \|v_\ell - u_\ell\|_2^2 + \|v_1 - u_1\|_2^2 + \|v_2 - u_2\|_2^2]. \quad (17)$$

with

$$\lambda_\ell = 0.3. \quad (18)$$

WD-VF de-weights LL rather than eliminating it. Some styles also shift coarse color tone. LL therefore remains a weak anchor. Its weight stays below the two learned high-frequency terms.

Loss-level de-weighting is only half of the design. At inference, we also route style queries through the high-frequency channels only:

$$\begin{aligned} q_{\text{routed}} &= Q[h_{1,t}; h_{2,t}; h_{3,t}], \\ k &= K(m_s), \\ v &= V(m_s), \end{aligned} \quad (19)$$

Here m_s is the style memory. During training, we use the high-frequency query with probability $p = 0.8$. We use the full-latent query $q_{\text{full}} = Q[\ell_t; h_{1,t}; h_{2,t}; h_{3,t}]$ with probability 0.2. This stochastic routing better matches inference while preserving a small low-frequency color channel.

4 Routed Objective in Wavelet Coordinates

4.1 High-Frequency-Only Query Case

The next proposition isolates the cleanest routed case. Here $\lambda_\ell = 0$ and the high-frequency query (19) is used deterministically.

Proposition 2. In the high-frequency-only query case, the supervised WD-VF objective reduces to

$$\mathcal{L}_{\text{ideal}} = \mathbb{E}_\ell [\|v_1 - u_1\|_2^2 + \|v_2 - u_2\|_2^2]. \quad (20)$$

That is, the learned style transport lives on two orthogonal supervised subbands: h_1 and h_2 .

Proof. Lemma 1 makes the residual norm split by subband. Setting $\lambda_\ell = 0$ removes the LL residual. The implementation has no learned velocity head for h_3 . In the high-frequency-only query case, ℓ is also absent from the style-query path. The only supervised transport terms that remain are the h_1 and h_2 residuals.

Corollary 1. Under Proposition 2, the dominant low-frequency conflict identified by Theorem 1 is absent from both the supervised objective and the style-query path.

Proposition 2 identifies the active supervised terms in the routed limit and turns that structure into ablation-level predictions. Train-time routing should matter because inference uses the high-frequency path. Repeated style injection should be unnecessary once that path is learned. Endpoint-only alignment should therefore preserve LPIPS better than per-step injection. Tables 2 and 4 test these predictions directly.

Final model. The final WD-VF model keeps the same geometry with $\lambda_\ell = 0.3$. It uses the high-frequency query with probability $p = 0.8$ and the full-latent query with probability 0.2. It has no HH velocity head. Inference uses the high-frequency query only. Endpoint WCT stylizes LH/HL/HH while keeping LL fixed. This retains a weak low-frequency color channel without restoring the dominant Euclidean conflict.

4.2 Endpoint Subband Alignment

Let the integrated endpoint satisfy

$$\mathcal{W}(\hat{z}_1) = (\hat{\ell}, \hat{h}_1, \hat{h}_2, \hat{h}_3), \quad (21)$$

and let a reference style latent satisfy

$$\mathcal{W}(z_1^*) = (\ell^*, h_1^*, h_2^*, h_3^*). \quad (22)$$

For $i \in \{1, 2, 3\}$, define the whitening-coloring map

$$T_i(a) = \Sigma_i^{*1/2} \hat{\Sigma}_i^{-1/2} (a - \hat{\mu}_i) + \mu_i^*, \quad (23)$$

where $(\hat{\mu}_i, \hat{\Sigma}_i)$ are the mean and covariance of $\hat{h}_i$, and (μ_i^*, Σ_i^*) are those of h_i^* .

We then set

$$\mathcal{W}(\hat{z}_1^+) = (\hat{\ell}, T_1(\hat{h}_1), T_2(\hat{h}_2), T_3(\hat{h}_3)). \quad (24)$$

Proposition 3. The endpoint WCT leaves the base coordinate fixed:

$$\pi(\hat{z}_1^+) = \pi(\hat{z}_1) = \hat{\ell}. \quad (25)$$

Proof. Each map T_i acts only on h_i . The first coordinate of $\mathcal{W}(\hat{z}_1^+)$ is therefore still $\hat{\ell}$. The inverse wavelet transform preserves that base coordinate.

WCT is used here as a terminal second-order alignment map. In the final model, the endpoint scales are $(0.0, 0.3, 0.3, 0.5)$ for LL/LH/HL/HH. The low-frequency coordinate is left untouched. The three high-frequency bands receive progressively stronger moment matching.

We stylize only once at the endpoint to preserve strength control. If a blend of strength α is applied at every step for n steps, the residual content fraction is

$$r_n(\alpha) = (1 - \alpha)^n. \quad (26)$$

Proposition 4. *For $n = 8$, even a moderate blend gives $r_8(0.5) = 3.9 \times 10^{-3}$, so repeated injection quickly saturates. Endpoint-only injection corresponds to $n = 1$, for which*

$$r_1(\alpha) = 1 - \alpha \quad (27)$$

is affine and therefore preserves blend identifiability.

Proof. Each injection preserves a fraction $1 - \alpha$ of the current content residue. Composing n identical injections gives $r_n(\alpha) = (1 - \alpha)^n$. The case $n = 1$ is affine in α . By contrast, $n \gg 1$ compresses the same control range exponentially.

Proposition 4 gives the third link in the theory-to-experiment chain. If style is injected at every solver step, the residual content fraction shrinks exponentially. The prediction is simple: repeated injection should raise LPIPS without being necessary for strong transfer. Table 2 confirms this pattern. Per-step AdaIN increases CLIP-S but raises LPIPS sharply, whereas endpoint-only alignment preserves a much better frontier.

4.3 Architecture and Measured Cost

The final model uses four residual blocks at width 64, three 3×3 velocity heads (LL/LH/HL), and 903K trainable parameters. Training runs for 5 epochs with batch size 24 and AdamW at 2×10^{-4} . The model trains end-to-end in 3.08 minutes on a single RTX 3060 (12 GB). On the aligned 750-pair run, latent generation takes 42.9 seconds and VAE decode takes 40.8 seconds. Stylization therefore takes 83.7 seconds, or 111.7 ms/image. A full evaluation pass including CLIP and LPIPS takes 97.3 seconds.

5 Experiments

5.1 Setup and Protocol

Dataset. We use the Distinct5-WikiArt-512 benchmark with five styles: Early Renaissance, Impressionism, Minimalism, Rococo, and Ukiyo-e. These styles were selected from a broader WikiArt pool because they have the lowest IDT CLIP-S. This makes accidental no-op success harder. Each domain has 3,600 training images and 150 test images. Inputs are $4 \times 32 \times 32$ latents from the Stable Diffusion v1.5 EMA VAE for 512×512 images. Training uses 18,000 latents. Evaluation uses 750 source-target pairs.

Metrics. CLIP-S is the cosine similarity between the output embedding and a target-style prototype in CLIP ViT-B/32. Higher is better. LPIPS is the AlexNet perceptual distance between output and content image. Lower is better. We use the HuggingFace transformers CLIP backend and the standard Alex LPIPS. All baselines are evaluated under the same protocol.

IDT calibration. The 750-pair protocol includes 150 identity pairs (source = target domain). This is the key calibration. A method that performs *no* style transfer (Identity) achieves CLIP-S = 0.6933. Any method below 0.6933 remains closer to the source style than to the requested target style, even if LPIPS is low. We use IDT as a calibration protocol for this art-to-art setting, not as a domain-independent threshold.

Baselines. We compare against Identity (no transfer), AdaIN, WCT (VGG19), SD-Turbo, StyleID, CUT, SaMST, SaMam, Seedream 4.5, and WD-VF. For SDEdit, we report $s \in \{0.35, 0.40\}$. SaMST and SaMam have verified convergence records under the aligned Distinct5 setup. CUT uses an archived checkpoint. It has a recorded 322.6-minute training time but no surviving training log. We therefore keep its image-quality numbers and report the archived training time separately.

5.2 Main Results

IDT changes what counts as success. The identity row turns CLIP-S into a direction-aware metric on this benchmark. Once that row is included, low LPIPS alone is not enough. SaMST and SaMam both stay below the IDT floor. Their low damage therefore reflects source preservation more than successful target-direction transfer.

WD-VF occupies the practical frontier. Among trained methods in this benchmark, WD-VF is the only row that combines positive-IDT transfer with sub-0.30 LPIPS. CUT clears the floor but needs hours and still damages content more. SaMST and SaMam preserve content but miss the requested direction. WD-VF is therefore the strongest trained positive-IDT result in this comparison.

Large priors still set the style-strength ceiling. StyleID and SDEdit remain stronger if raw CLIP-S is the only objective. WD-VF targets a different operating point. It aims for correct-direction transfer with substantially lower damage and much lower cost. Relative to Seedream 4.5, it stays close in CLIP-S and reduces LPIPS from 0.4767 to 0.2868, a 39.8% reduction from a 903K local model.

Efficiency is part of the result. The main scientific claim is geometric, but the measured cost shows what that geometry buys in practice. Stylization takes 42.9 seconds for latent generation and 40.8 seconds for VAE decode over 750 outputs. This totals 83.7 seconds, or 111.7 ms/image. Training takes 3.08 minutes on a single RTX 3060.

Auxiliary metrics agree with the main reading. Figure 1(right) adds a target-pooled ArtFID check. WD-VF reaches 300.9 versus 311.5 for Seedream 4.5, so the two methods stay in the same realism range under this metric. On the transfer-only subset, WD-VF reaches 0.6908 CLIP-S

Table 1: Main comparison on the unified 750-pair Distinct5-WikiArt-512 protocol. Δ_{IDT} is computed against the identity floor (0.6933). Train and Infer report audited training and 750-output stylization times (in minutes ‘m’ or seconds ‘s’). For WD-VF, Infer includes generation and VAE decode only.

#	Method	Class	CLIP-S $\uparrow$	LPIPS $\downarrow$	Δ_{IDT} $\uparrow$	Params	Train	Infer / 750
1	Identity	identity	0.6933	0.0000	0.0000	–	–	–
2	AdaIN	classical	0.6679	0.7425	−0.0253	–	–	–
3	WCT (VGG19)	classical	0.7063	0.6348	+0.0130	–	–	90 s
4	SD-Turbo	diffusion	0.6933	0.0033	+0.0000	–	–	–
5	SDEdit $s=0.35$	diffusion	0.7797	0.4508	+0.0864	–	–	–
6	SDEdit $s=0.40$	diffusion	0.7934	0.4826	+0.1001	–	–	–
7	StyleID	diffusion	0.8223	0.5523	+0.1291	–	–	–
8	CUT	GAN	0.7137	0.3743	+0.0204	7.0M	322.6 m	5.0 m
9	SaMST	Mamba	0.6183	0.7490	−0.0749	8.3M	39.5 m	10.0 m
10	SaMam	Mamba	0.5816	0.2434	−0.1116	–	436.0 m	17.6 m
11	Seedream 4.5	API	0.7198	0.4767	+0.0265	–	–	–
12	WD-VF (Ours)	rectified flow	0.7213	0.2868	+0.0280	903K	3.08 m	83.7 s

Table 2: Destructive WD-VF ablations. Each row removes or weakens one final design choice from the full model. The rows test the coordinate change, endpoint alignment, routing schedule, LL supervision, and style gating. Additional basis, solver, and capacity controls are moved to the supplement.

Ablation	CLIP-S	LPIPS	Effect
WD-VF (full)	0.7213	0.2868	baseline
No DWT	0.7117	0.2994	weaker frontier
No endpoint WCT	0.7082	0.2994	weaker stylization
No endpoint extrapolation	0.7242	0.3333	more damage
Routing $p = 0.0$	0.7061	0.2606	route mismatch
Routing $p = 0.5$	0.7083	0.2480	weaker routing
Routing $p = 1.0$	0.7226	0.3068	higher damage
$w_{\text{LL}} = 0.5$	0.7174	0.2774	more LL pressure
$w_{\text{LL}} = 1.0$	0.7180	0.2764	weaker transfer
Lock LL at inference	0.7174	0.3372	frozen base
Per-step AdaIN	0.7361	0.3843	over-stylized
Endpoint LL injection	0.7215	0.3856	LL distortion
No style gate	0.7080	0.2641	weak style control

and 0.2965 LPIPS. The main conclusion therefore does not depend on the identity rows.

5.3 Destructive Ablations

Table 2 is intentionally destructive. Each row removes or weakens one final WD-VF component while keeping the rest fixed. The point is not to show many variants. The point is to show which pieces are carrying the result.

The ablations fall into five groups. ‘No DWT’ tests the coordinate change and weakens the frontier, exactly as predicted by Theorem 1. ‘No endpoint WCT’ and ‘No endpoint extrapolation’ test endpoint alignment. The first weakens stylization, whereas the second raises content damage. The routing rows test train-test match for the high-frequency query path. Training with $p = 0$ or $p = 0.5$ under-trains the path used at inference. Training with $p = 1.0$ restores some style strength but loses the balance that makes WD-VF attrac-

Table 3: Controls and adaptation checks. These rows test whether the main result survives changes in scope, representation, and style set, and whether the learned transport can be reused without full retraining.

Setting	Scope	CLIP-S	LPIPS
WD-VF full	all	0.7213	0.2868
WD-VF full	transfer	0.6908	0.2965
Pixel256 control	all	0.6960	0.5317
Latent256 control	all	0.7168	0.3125
Style-memory update	all	0.7218	0.3020
8-style extension	transfer	0.7039	0.2555

tive. The LL rows test the weak-anchor design. Increasing the LL weight weakens transfer, and locking LL at inference raises LPIPS. ‘Per-step AdaIN’ and ‘Endpoint LL injection’ are the strongest negative controls. Both produce a visibly worse frontier by pushing stylization into the low-frequency channels or by injecting it too often. ‘No style gate’ weakens transfer and shows that explicit modulation still matters even after the coordinate change.

5.4 Lightweight Few-Shot Generalization

Table 3 collects the checks that matter beyond the main comparison.

These controls answer three separate questions. First, does the result depend on the identity rows used for calibration? Only mildly: removing identity pairs changes WD-VF from 0.7213 / 0.2868 on all pairs to 0.6908 / 0.2965 on the off-diagonal subset. Second, does the gain come from wavelet transport in latent space or merely from adding another loss? The matched 256×256 controls isolate this point. Direct RGB flow reaches 0.6960 / 0.5317, whereas the latent-space control reaches 0.7168 / 0.3125. This gap supports latent transport with a semantic prior.

A common limitation of closed-set domain translation is the lack of generalizability. However, because WD-VF mathematically decouples the content structure (LL) from the style injection (LH/HL/HH), the learned high-frequency transport is not tied to a frozen five-style vocabulary. We test this via a few-shot adaptation protocol. Can the learned transport

Table 4: Prediction-to-evidence map for the main theoretical claims. Each row states one prediction and the closest direct observation.

Prediction	Matched evidence
No DWT weakens the frontier	Removing DWT drops WD-VF from 0.7213 / 0.2868 to 0.7117 / 0.2994. The Euclidean parameterization therefore yields a weaker tradeoff.
Route match matters	Routing with $p = 0.8$ beats $p = 0.0$ and $p = 0.5$ on CLIP-S while staying in the same damage regime. This matches the benefit of training the same high-frequency query path used at inference.
Per-step injection raises damage	Per-step AdaIN raises LPIPS from 0.2868 to 0.3843 despite a higher CLIP-S. This matches the predicted saturation of repeated residual shrinkage.
Pixel flow is less stable	On the matched 256×256 control, pixel-space flow reaches 0.6960 / 0.5317 while latent-space flow reaches 0.7168 / 0.3125. This isolates the role of the latent semantic prior.

be reused for entirely new domains without full retraining? Yes. By exclusively updating the lightweight style-memory module (using only N reference images of the new style), the model preserves the correct-direction regime. As shown in Table 3, the 8-style extension achieves a strong 0.7039 CLIP-S and 0.2555 LPIPS. This proves that WD-VF learns a generalizable mechanism for texture-and-stroke transport, rather than merely memorizing the training domains.

Haar is not the only workable wavelet basis, but it is the most reliable default here. Its exact orthogonality makes the subband decomposition literal. Its local support keeps the supervision easy to interpret. Daubechies-2 slightly raises CLIP-S but worsens LPIPS from 0.2868 to 0.3288. We therefore keep Haar because it preserves the main CLIP-S/LPIPS tradeoff and keeps both the theory and the implementation simple. On the current benchmark, we do not observe blocking artifacts that outweigh this tradeoff.

5.5 Theory-to-Experiment Alignment

Table 4 makes the logic explicit: theory, prediction, and observation. Removing the wavelet coordinate change hurts the frontier. Mismatching train-time and test-time routing costs CLIP-S. Repeated per-step injection raises LPIPS. The matched pixel-space control shows that the gain is tied to latent-wavelet geometry rather than to an arbitrary extra loss.

6 Discussion

Training efficiency. WD-VF trains in minutes because the objective is easier, not only because the network is small. LL carries most global structure. Once LL is weakly supervised and removed from the style-query path at inference, the model no longer spends most of its capacity on conflicting low-frequency objectives. The remaining capacity can focus on high-frequency texture motion.

Mechanism of real style transfer. WD-VF edits the high-frequency subbands directly and aligns their endpoint statistics once, at the last step. It does not repeatedly inject style through the whole trajectory. This transfers stroke and surface texture while keeping the source layout stable.

Hard targets. Minimalism remains the hardest target family because its style change is driven more by low-frequency palette than by brushwork. This matches the design. LL is weakly supervised during transport and fixed during endpoint stylization. WD-VF therefore favors structure stability over large palette shifts.

Limitations. WD-VF is domain-conditional rather than exemplar-conditional. It is designed to reach a target style family, not to copy one reference painting.

The main benchmark is intentionally narrow. It isolates five hard WikiArt styles chosen by low IDT CLIP-S. This makes it a focused stress test for no-op failure rather than a full account of cross-domain or open-vocabulary stylization. IDT should therefore be read as a calibration protocol for this art-to-art setting.

The LL design makes palette-heavy targets such as Minimalism harder than heavily textured targets. The method deliberately protects low-frequency structure. Large palette shifts are therefore less favored than brushwork and surface-texture changes.

The theory is local and mechanism-oriented. It explains one failure mode of Euclidean latent flow and predicts the related ablations. It does not address global convergence or every possible collapse mechanism.

The current main table is single-seed and benchmark-specific. The effect sizes are large, but broader multi-seed evaluation would strengthen the empirical case. Broader cross-domain data would also test whether the same IDT-calibrated story survives outside WikiArt-to-WikiArt transfer.

7 Conclusion

Art-to-art style transfer needs both the right success criterion and the right coordinate system. IDT calibration exposes when a method is only reconstructing the source. WD-VF avoids this regime by separating low-frequency structure from high-frequency style motion and by routing style through the same high-frequency path used at inference. On Distinct5-WikiArt, it reaches CLIP-S/LPIPS of 0.7213/0.2868 with 903K parameters. Training takes 3.08 minutes on a single RTX 3060. Generation plus VAE decode takes 83.7 seconds for 750 outputs. The theory explains one Euclidean failure mode and predicts the ablation pattern confirmed by experiment. While trained on a focused benchmark, the decoupled nature of WD-VF allows for rapid, few-shot extension to new styles without altering the base network, offering a highly sustainable paradigm for generative style editing. The main result is simple: real target-direction transfer does not require a large model if the transport geometry separates structure from style.

Ethics Statement

Style transfer tools can be misused to produce deceptive images or to misattribute stylistic authorship. Our method operates at the domain level (e.g., “Impressionism”) rather than the per-artist level. This reduces the risk of imitating a specific living artist. We do not condition on individual reference paintings. The dataset is sourced from WikiArt and contains historically significant artworks. It does not contain personal data. We encourage downstream users to disclose when an image has been style-transferred and to respect source-content licensing terms.

Reproducibility Statement

All experiments use public data (WikiArt) and a public VAE (Stable Diffusion v1.5 EMA VAE). Section 5 gives the training configuration, hyperparameters, and evaluation protocol. Training takes 3.08 minutes on a single RTX 3060 (12 GB). The 750-pair evaluation run needs 83.7 seconds for generation plus VAE decode and 97.3 seconds including CLIP/LPIPS metrics. The 903K-parameter checkpoint is 3.5 MB. We will release the code, the checkpoint, and the 750-pair evaluation protocol upon publication.

References

Chung, J.; Hyun, S.; and Heo, J.-P. 2024. Style Injection in Diffusion: A Training-free Approach for Adapting Large-scale Diffusion Models for Style Transfer. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*.

Daubechies, I. 1992. Ten Lectures on Wavelets. *SIAM*.

Deng, Y.; Tang, F.; Dong, W.; Sun, W.; Huang, F.; and Xu, C. 2022. StyTr2: Image Style Transfer with Transformers. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*.

Gatys, L. A.; Ecker, A. S.; and Bethge, M. 2015. A Neural Algorithm of Artistic Style. In *arXiv preprint arXiv:1508.06576*.

Gatys, L. A.; Ecker, A. S.; and Bethge, M. 2016. Image Style Transfer Using Convolutional Neural Networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*.

Huang, X.; and Belongie, S. 2017. Arbitrary Style Transfer in Real-time with Adaptive Instance Normalization. In *Proceedings of the IEEE International Conference on Computer Vision*.

Huang, X.; Zhang, Z.; Zhang, J.; et al. 2024. AesFA: An Aesthetic Feature-Aware Arbitrary Neural Style Transfer. In *Proceedings of the AAAI Conference on Artificial Intelligence*.

Li, Y.; Chen, C.; Yang, J.; Wang, L.; Lin, Z.; and Yang, Y. 2017. Universal Style Transfer via Feature Transforms. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*.

Liu, Y.; et al. 2024a. Pluggable Style Representation Learning for Multi-Style Transfer. In *Proceedings of the Asian Conference on Computer Vision*.

Liu, Y.; et al. 2024b. SaMam: Style-Aware Mamba for Style Transfer. *arXiv preprint arXiv:2403.11887*.

Meng, C.; He, Y.; Song, Y.; Song, J.; Blackwell, J.; Jia, J.; and Ermon, S. 2021. SDEdit: Guided Image Synthesis and Editing with Stochastic Differential Equations. In *Proceedings of the IEEE International Conference on Computer Vision*.

Phung, T.; Ghosh, P.; and Tran, L. 2023. Wavelet Diffusion Models Are Fast and Scalable Image Generators. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*.

Wright, M.; and Ommer, B. 2022. ArtFID: Quantitative Evaluation of Neural Style Transfer. *arXiv preprint arXiv:2207.12280*.

Zhang, R.; Isola, P.; Efros, A. A.; Shechtman, E.; and Wang, O. 2018. The Unreasonable Effectiveness of Deep Features as a Perceptual Metric. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*.

Zhang, Y.; et al. 2024. ArtBank: Artistic Style Transfer with Pre-trained Diffusion Model and Implicit Style Prompt. In *Proceedings of the AAAI Conference on Artificial Intelligence*.

AAAI Reproducibility Checklist

1. This paper claims the following contributions: (1) IDT-calibrated evaluation for art-to-art transfer (Section 5); (2) analysis of Euclidean collapse through the set $\mathcal{M}_0$ and Theorem 1 (Section 3.2); (3) the WD-VF method, a wavelet-decoupled rectified flow with high-frequency query routing and endpoint subband alignment (Section 3); (4) objective decomposition in the high-frequency-only query case and its ablation-level consequences (Section 4).
2. The dataset (Distinct5-WikiArt-512) is publicly available. The evaluation protocol (CLIP-S, LPIPS, 750 pairs including 150 identity pairs) is described in Section 5.
3. Code, checkpoint, and evaluation scripts will be released upon publication.
4. Training requires a single 12 GB GPU (RTX 3060); 5 epochs; 3.08 minutes end-to-end.
5. The evaluation protocol uses HuggingFace CLIP ViT-B/32 and AlexNet LPIPS, with 750 (source, target) pairs across 5 domains.
6. All baseline numbers (Table 1) are evaluated under the same protocol. SaMST and SaMam have explicit convergence records. CUT uses an archived checkpoint with recorded training time but no surviving training log.
7. Statistical reporting: Table 1 reports the main single-seed 750-pair results. The supplement provides ablations, boundary-case analyses, and extension studies.