

Latent Bridge Matching: Efficient OT-Coupled Flow Matching for Domain-Level Artistic Transfer

Anonymous submission

Abstract

We propose Latent Bridge Matching, an efficient framework for domain-level artistic transfer that learns a time-conditioned velocity field in compact VAE latent space. Rather than optimizing per-image or sampling a large diffusion model, our method constructs OT-coupled content-style latent pairs using a sliced-Wasserstein transport cost, trains a latent flow-matching objective at randomly sampled bridge times, and regularizes the integrated endpoint with terminal distribution matching. This separates two forces: the OT/SWD terminal objective drives the output toward the target artistic domain, while the latent flow and kinetic regularization constrain the path to preserve content structure. Under a strict 750-output protocol over five domains, our 3.9M-parameter model reaches competitive CLIP-style performance, improves LPIPS and EC over SaMST, S2WAT, AdaIN, CAST, and StyleID, and trains substantially faster than heavier feed-forward baselines. The results suggest that OT-coupled latent flow matching provides a clean, reproducible efficiency-quality trade-off for domain-level artistic transfer.

Introduction

Artistic style transfer has long served as a testbed for separating image content from visual appearance. Classical neural style transfer showed that convolutional features and Gram statistics can synthesize compelling artistic textures [6], while feed-forward arbitrary style transfer made real-time deployment practical through normalization-based style modulation [1]. More recent attention, transformer, training-free, and diffusion-based systems have improved flexibility and style fidelity [1, 4, 9, 1, 7, 2]. However, practical multi-style artistic transfer still exposes a persistent three-way tension: strong target-style expression, faithful content preservation, and low computational cost.

This tension is especially visible in domain-level multi-style transfer. Efficient multi-style systems can support many styles and fast inference, but the output may become washed, over-smoothed, or contaminated by structured grain artifacts. Diffusion-based or training-free methods can produce strong style semantics, but may be expensive or suffer semantic drift

under content-preserving evaluation. Under our strict 750-output protocol, StyleID attains the highest raw CLIP-style score but exhibits severe content collapse, while SaMST achieves strong raw style and content scores yet produces visible muddy and grain-like texture artifacts. These observations suggest that standard style/content metrics alone are insufficient. A competitive system must also be evaluated for perceptual quality and artifact-sensitive texture realism.

We propose Latent Bridge Matching (LBM), which models stylization as OT-coupled latent flow matching. Instead of directly predicting RGB images, the method operates on compact VAE latents with three integrated components: (i) a mini-batch entropic OT coupling that matches unpaired content and target-style latents via sliced-Wasserstein cost; (ii) a time-conditioned velocity field trained with flow-matching supervision at randomly sampled bridge times; and (iii) a terminal SWD regularizer that pulls the Euler-integrated endpoint toward the target style distribution. A learnable style spatial prior and semantic cross-attention provide efficient domain-level conditioning without per-image optimization or diffusion sampling.

Our contributions are:

- We formulate domain-level artistic transfer as OT-coupled latent flow matching, where unpaired content and target-style latents are matched with a sliced-Wasserstein transport cost and used to supervise a time-conditioned velocity field.
- We introduce a lightweight domain-conditioned latent backbone with learned style embeddings, spatial priors, and semantic routing, enabling style transfer without per-image optimization, reference-image encoding, or diffusion sampling.
- We combine flow-matching supervision with terminal SWD endpoint regularization, explicitly separating path/content preservation from target-domain distribution alignment.
- We evaluate under a strict 750-output protocol and show that our method lies on a favorable efficiency-quality Pareto frontier: 3.9M parameters, competitive CLIP-style, lower LPIPS, best EC among compared methods, and substantially lower training

cost.

Related Work

Classical and arbitrary neural style transfer. Neural style transfer began with optimization-based feature reconstruction and Gram-matrix style losses [6]. AdaIN later enabled real-time arbitrary transfer by aligning channel-wise feature statistics [1]. Attention-based methods such as SANet improved local style matching through style-attentional feature aggregation [1]. Recent aesthetic and pattern-aware methods, including AesPA-Net and AesFA, further emphasize pattern repeatability, aesthetic features, and high-quality arbitrary stylization [9, 1]. These methods demonstrate the value of feature statistics and local matching, but they generally do not explicitly model a style-conditioned transport path with a terminal distribution constraint.

Transformers and multi-style transfer. Transformer-based style transfer methods such as StyTr2 and S2WAT use attention mechanisms to improve content-style interaction and long-range feature modeling [4, 1]. Multi-style approaches amortize many styles in one compact system; SaMST is a recent representative with pluggable style representations, low parameter count, and fast inference [1]. These systems are directly related to our goal of efficient multi-style transfer. In contrast to representation-only style injection, our method places semantic cross-attention inside a latent bridge objective, so terminal distribution matching and kinetic regularization explicitly shape the style-content trade-off.

Training-free and diffusion methods. Training-free and diffusion-based approaches provide a different route to flexible style transfer. CAST uses VQ autoencoders for content-adaptive training-free style transfer [7], while StyleID injects style into large-scale diffusion models without training [2]. These methods can be visually strong, but diffusion sampling and semantic drift can limit their suitability for fast domain-level transfer. Our work targets a complementary efficiency regime: compact latent integration with style-id conditioning rather than iterative diffusion sampling or per-style optimization.

Distribution matching and perceptual evaluation. Our objective is motivated by optimal transport and distribution matching. Sinkhorn distances and computational optimal transport provide efficient tools for approximate transport problems [3, 1], while the Schrödinger problem connects entropic transport to stochastic path measures [1]. We use SWD as a lightweight terminal distribution matching objective. For evaluation, we combine perceptual and distributional measures including LPIPS [2], FID [8], KID [1], DISTS [5], MUSIQ [1], MANIQA [1], and style-transfer-specific ArtFID [1]. This broad metric set is important because raw CLIP-style and structural scores may fail to penalize grain-like artifacts.

Method

Overview

We model domain-level artistic transfer as OT-coupled latent flow matching. Let $z_0 \in \mathbb{R}^{4 \times 32 \times 32}$ denote the VAE latent of a content image and let s denote a target style domain identifier. The method has three components: (i) an optimal-transport coupling that matches unpaired content and target-style latents within each minibatch; (ii) a time-conditioned velocity field $v_\theta(z_t, t, s)$ trained with flow-matching supervision at randomly sampled bridge times; and (iii) a terminal sliced-Wasserstein regularizer that pulls the integrated endpoint toward the target style distribution. At inference, the learned vector field is integrated with explicit Euler steps and decoded by the VAE decoder.

Unpaired latent-domain sampling and OT coupling

We train directly on precomputed VAE latents without paired supervision. Each minibatch samples content latents $\{z_c\}$ from all source domains and target-style latents $\{z_s\}$ from the requested target domain. Identity pairs are naturally included through uniform target-domain sampling, which stabilizes the learned transport around same-domain mappings.

Rather than using random target assignments, we construct a mini-batch transport coupling using the sliced-Wasserstein cost. For a batch of content latents $\{z_c^{(i)}\}$ and target-style latents $\{z_s^{(j)}\}$, the SWD cost matrix is

$$C_{ij} = \text{SWD}(z_c^{(i)}, z_s^{(j)}). \quad (1)$$

An entropic optimal-transport plan is then obtained via the Sinkhorn algorithm:

$$\pi^* = \arg \min_{\pi} \sum_{ij} C_{ij} \pi_{ij} + \epsilon H(\pi), \quad (2)$$

where $H(\pi)$ is the entropy regularizer. Each content latent is matched to its coupled target $\tilde{z}_1 \sim \pi^*(z_s|z_c)$. This coupling provides style-domain endpoints without requiring spatially aligned targets.

Time-conditioned latent flow matching

The core of the method is a time-conditioned velocity field $v_\theta(z_t, t, s)$ that transports latents along a stochastic bridge. At a randomly sampled time $t \sim \mathcal{U}(0, 1)$, the bridge state is

$$z_t = (1-t)z_0 + t\tilde{z}_1 + \sigma\sqrt{t(1-t)}\epsilon, \quad \epsilon \sim \mathcal{N}(0, I), \quad (3)$$

and the target velocity is the derivative of the transport:

$$u_t = \tilde{z}_1 - z_0 + \sigma \frac{1-2t}{2\sqrt{t(1-t)}} \epsilon. \quad (4)$$

The flow-matching objective supervises the network to predict this velocity:

$$\mathcal{L}_{\text{FM}} = \mathbb{E}_{z_0, \tilde{z}_1, t, \epsilon} \|v_\theta(z_t, t, s) - u_t\|^2. \quad (5)$$

This formulation is equivalent to learning a denoising score along the bridge, and it naturally separates path content preservation from endpoint style matching.

Architecture. The time-conditioned backbone uses sinusoidal time embeddings added to the style code. The implementation is a U-Net-like latent trunk with four semantic cross-attention body blocks at 16×16 resolution. The target style is represented by a learnable spatial prior $P_s \in \mathbb{R}^{128 \times 16 \times 16}$. Semantic cross-attention injects the style prior into content features, while skip connections preserve spatial structure.

Terminal SWD regularization

In addition to flow-matching supervision, a terminal loss regularizes the Euler-integrated endpoint toward the target style distribution. Let $\Phi_\theta(z_0, s)$ denote the endpoint obtained by integrating v_θ from z_0 with style s . The terminal loss is

$$\mathcal{L}_{\text{term}} = \text{SWD}(\Phi_\theta(z_0, s), \tilde{z}_1), \quad (6)$$

where the SWD compares generated endpoint patches with coupled target patches using 64 projections over patch sizes $\{3, 5, 7, 15\}$.

The complete training objective combines flow matching, terminal distribution matching, and kinetic path regularization:

$$\mathcal{L} = \mathcal{L}_{\text{FM}} + \lambda_{\text{term}} \mathcal{L}_{\text{term}} + \lambda_{\text{kin}} \mathcal{L}_{\text{kin}}, \quad (7)$$

where $\mathcal{L}_{\text{kin}} = \mathbb{E} \|v_\theta\|_2^2$ penalizes excessive latent displacement. The primary configuration uses $\lambda_{\text{term}} = 20.0$ and $\lambda_{\text{kin}} = 1.0$. Color, cycle, NCE, and repulsive losses are disabled in the main setting. This three-part loss separates the forces that drive style from those that preserve content.

Bounded theoretical claim. The method is not an exact Schrodinger Bridge solver. It does not estimate forward and backward stochastic scores nor optimize the full entropic dual objective. The bridge perspective is used as a tractable training principle: a flow-matching path constraint, a terminal distribution regularizer, and a kinetic penalty. This bounded interpretation is reflected empirically—removing the terminal constraint weakens style, while removing the path penalty destroys content.

More concretely, let $\mathcal{P}_r(\hat{z}_1)$ be the set of latent patches of size r extracted from the generated endpoint, and let $\mathcal{P}_r(z_s)$ be the corresponding patch set from target-style samples. For projection directions $\{\theta_m\}_{m=1}^M$, the empirical terminal objective is

$$\mathcal{L}_{\text{swd}} = \sum_{r \in \mathcal{R}} \frac{1}{M} \sum_{m=1}^M \left\| \text{sort}(\mathcal{P}_r(\hat{z}_1) \theta_m) - \text{sort}(\mathcal{P}_r(z_s) \theta_m) \right\|_2^2. \quad (8)$$

Here $\mathcal{R} = \{3, 5, 7, 15\}$ in the primary configuration. The loss therefore asks whether the generated endpoint has target-style patch statistics, not whether it matches any particular artwork spatially.




Figure 1: Overview of the latent transport model.

Experiments

Protocol and metrics

We evaluate under a strict 750-output protocol: five source domains times five target domains times 30 source images. The five domains are photo, Hayao, Monet, Van Gogh, and Cezanne; Ukiyo-e is not used in the present protocol. We report CLIP-style, CLIP-content, LPIPS-content, and the composite score

$$\text{EC} = \text{CLIP-style} \cdot (1 - \text{LPIPS}). \quad (9)$$

For artifact-sensitive analysis we also report MUSIQ, MANIQA, DISTS-content, high-frequency patch KID, FFT slope error, and shallow Gram micro statistics. We run the strict-750 protocol end-to-end for Ours, SaMST, StyleID, S2WAT, and AdaIN variants.

Metric taxonomy. We use four metric groups. Raw semantic metrics, CLIP-style and CLIP-content, measure whether an output is recognized as stylistically aligned and semantically related to the input. Perceptual content metrics, LPIPS and DISTS-content, measure visual structure preservation. Distributional style metrics, including KID, Gram statistics, and high-frequency patch KID, evaluate whether generated texture statistics resemble target-style images. No-reference image-quality metrics, MUSIQ and MANIQA, help detect dirty or over-processed outputs. We use EC only as an internal Pareto diagnostic for sweep selection; all claims are checked against the individual metrics.




Figure 2: Strict-750 style-content trade-off. Ours is close to SamST in raw CLIP-style while retaining a better LPIPS-based trade-off than several baselines; StyleID reaches high style but collapses content.

Main comparison

Table 5 reveals a clear style-content-efficiency trade-off. StyleID obtains the highest raw CLIP-style score, but its LPIPS of 0.750 indicates severe content drift and its diffusion-based inference is over $15\times$ slower. SamST is the closest efficient competitor and slightly exceeds our CLIP-style by 0.003; however, Ours achieves lower LPIPS (0.451 vs 0.466), higher EC (0.393 vs 0.384), fewer parameters (3.9M vs 6.0M), and substantially lower training cost (310 s vs 6769 s). S2WAT reaches a similar CLIP-style but has much worse LPIPS (0.526) and a $16.7\times$ larger model. AdaIN remains the fastest method, but its LPIPS and EC are far behind (0.630, 0.264). CAST and diffusion-based StyleID are constrained by content drift or high inference cost. These results support our main claim: Ours does not maximize a single metric alone, but provides a favorable Pareto point for efficient domain-level artistic transfer.

Artifact-sensitive comparison to SamST

SamST is a strong baseline under raw style and structure metrics, but visual inspection reveals muddy and grain-like artifacts. We therefore examine artifact-sensitive diagnostics in Table 6. Ours improves MUSIQ by 13.1109, MANIQA by 0.0918, DISTS-content by 0.0466, HF-Patch-KID by 2.5904, and FFT slope error by 0.5063. These results support the interpretation that SamST preserves low-frequency structure but produces less realistic local texture statistics.

This diagnostic block addresses a specific evaluation pitfall. Metrics such as CLIP-content, SSIM, and edge overlap tend to reward low-frequency structure preservation. A grainy output can therefore score well if object boundaries remain visible, even when the surface texture is visually dirty. The added diagnostics target a different failure mode: whether local high-frequency patches, no-reference perceptual quality, and frequency slopes resemble real artistic texture rather than spa-




Figure 3: Destructive ablations. Removing terminal SWD improves content but weakens style, while removing kinetic regularization increases style at the cost of severe content degradation.

tially incoherent dithering. These metrics should not replace human evaluation, but they provide objective evidence for the artifact mode observed in visual inspection.

Destructive ablations

The ablation results in Table 7 expose the mechanism of the model. Removing terminal SWD gives excellent content preservation but weak style, showing that terminal distribution matching is the primary style driver. Removing kinetic regularization preserves style but collapses content, confirming that the path-energy penalty is essential. Strong color loss damages both content and distributional quality, so naive color matching is not used in the mainline.

Weight sweep and theory switches

The 40-run category-weight sweep trained 20 sampling recipes with $K \in \{1, 2\}$ for eight epochs each, evaluating every epoch for 320 total checkpoints. The best EC result was K2 balanced default at epoch 3 with CLIP-style 0.6980, CLIP-content 0.8727, LPIPS 0.3777, and EC 0.4343. The best raw-style result in the sweep was K1 balanced default at epoch 8 with CLIP-style 0.7161, CLIP-content 0.7984, LPIPS 0.4605, and EC 0.3863. The sweep showed that more complex category weighting was not consistently better than balanced sampling.

We also evaluated optional theory switches. Sinkhorn-normalized semantic routing and entropy-gated kinetic regularization improved EC/content but reduced raw CLIP-style. The best theory-switch EC was entropy-gated kinetic with weight 5.0 at epoch 1, giving CLIP-style 0.6916, CLIP-content 0.8804, LPIPS 0.3684, and EC 0.4368. Color transport increased raw style slightly but harmed LPIPS and content, and was therefore not selected as the mainline.

Table 1: Strict-750 comparison with efficiency. $\uparrow$ = higher is better, $\downarrow$ = lower is better. Inference time end-to-end for 750 images. Training times marked $\dagger$ are extrapolated.

Method	Params	CLIP-S $\uparrow$	LPIPS $\downarrow$	EC $\uparrow$	Infer	s/img	Train sec
Ours ep7	3.9 M	0.716	0.451	0.393	40.0 s	0.053	309.9
SaMST	6.0 M	0.719	0.466	0.384	39.8 s	0.053	6768.7 $\dagger$
S2WAT	65 M	0.714	0.526	0.338	45.1 s	0.060	10600 $\dagger$
AdaIN v32k	$\sim$ 5 M	0.713	0.630	0.264	9.3 s	0.012	9220.4
StyleID	SD-based	0.760	0.750	0.190	603.3 s	0.804	training-free
CAST	7.0 M	0.665	0.726	0.182	75.5 s	0.101	training-free

$\dagger$ Extrapolated.

Table 2: Artifact-sensitive comparison between Ours and SaMST.

Metric	Ours e7	SaMST
MUSIQ $\uparrow$	49.2059	36.0950
MANIQA $\uparrow$	0.4057	0.3139
DISTS-content $\downarrow$	0.2477	0.2943
HF-Patch-KID $\downarrow$	4.1694	6.7598
FFT slope error $\downarrow$	0.5473	1.0536
Gram micro $\downarrow$	0.0798	0.0947

Table 3: Destructive ablations under the strict-750 protocol.

Variant	CLIP-S	CLIP-C	LPIPS	EC
D0 full	0.7014	0.8022	0.4593	0.3791
D1 w/o terminal SWD	0.6708	0.8989	0.3490	0.4368
D2 w/o kinetic	0.7159	0.6624	0.6375	0.2596
D8 strong color loss	0.6923	0.6629	0.5675	0.2994

Step-count and λ sensitivity

The additional step-count sweep evaluates 1, 4, 8, 12, and 16 Euler steps for the selected epoch-7 checkpoint. The main observation is that the quality curve is almost flat over this range. The all-pairs CLIP-style score varies only from 0.7159 to 0.7162, while LPIPS remains tightly concentrated around 0.4514. The photo-to-art operating point is similarly stable: CLIP-style changes only from 0.6712 to 0.6715 and LPIPS remains near 0.4881. This means the 12-step default is not a brittle tuned optimum. In fact, four steps already recover nearly identical quality, while the measured end-to-end strict-750 throughput increases from 8.43 images/s at one step to 9.54 images/s at four steps and remains around 9.34–9.50 images/s for 8–16 steps. We therefore keep 12 steps as a conservative default, but the sweep indicates that lower-step deployment is viable when marginal latency reduction matters.

We also scanned a 3×3 grid over $\lambda_{\text{kin}} \in \{0.5, 1.0, 2.0\}$




Figure 4: Summary of the 40-run category-weight sweep. K2 balanced default gives the best EC, whereas K1 balanced default gives the best raw CLIP-style among evaluated checkpoints.

and $\lambda_{\text{swd}} \in \{10, 20, 30\}$ for eight epochs each. The resulting surface is smooth and interpretable rather than erratic. Increasing kinetic weight while reducing terminal pressure systematically improves content preservation: at epoch 8, the setting (2.0, 10) yields CLIP-content 0.8734, LPIPS 0.3720, and EC 0.4373, compared with the default (1.0, 20) setting at CLIP-content 0.8124, LPIPS 0.4477, and EC 0.3937. Conversely, weaker kinetic weight or stronger terminal weight pushes the model toward more aggressive but less content-faithful stylization; for example, (0.5, 30) drops CLIP-content to 0.6926 and raises LPIPS to 0.5538. The best transfer-EC checkpoints in this grid occur consistently around epoch 3, again showing that the protocol exposes a broad style-content frontier rather than a single fragile endpoint.

Resource use and reproducibility

End-to-end inference and training times are listed alongside the main metrics in Table 5. All timings were measured on an NVIDIA RTX 4070 Laptop GPU (8 GB VRAM) with the same strict-750 protocol. Training times marked $\dagger$ are extrapolated from a single-epoch probe. StyleID is training-free by design.

All reported strict-750 evaluations use the same

domain protocol and preprocessing pipeline. We also evaluated every epoch in the 40-run sweep, yielding 320 checkpoint evaluations. Final-epoch selection alone can hide style-content trade-offs: K2 balanced default achieved the best EC at epoch 3, whereas K1 balanced default achieved the best raw style at epoch 8. To improve reproducibility, each training run stores the resolved runtime configuration, source-file snapshot, checkpoint, and per-epoch CSV logs. Python, NumPy, PyTorch, and DataLoader worker seeds are fixed per run.

Discussion and Limitations

The experiments support a specific and conservative conclusion: the proposed method lies on a favorable efficiency-quality Pareto frontier for domain-level artistic transfer. Ours does not achieve the single highest raw CLIP-style—StyleID attains 0.760 at the cost of content collapse, and SaMST is 0.003 ahead with heavier parameters and much longer training. However, Ours combines competitive style quality (CLIP-style 0.716) with the lowest LPIPS among high-style methods, the best EC score (0.393), the smallest model (3.9M), and the shortest measured training time (310 s). This joint efficiency-quality profile is the paper’s main finding.

The current study still has several limitations. First, the style representation is domain-level rather than reference-image-specific; this makes inference efficient but limits the ability to reproduce unique single-image brushstroke patterns or resolve cross-domain prototype conflicts. Second, Euler integration is deliberately simple, so discretization error can accumulate when the latent path is extrapolated to more aggressive settings or higher resolutions. Third, the SWD term uses heuristic multi-scale patch projections rather than an exact transport solver, so the bridge interpretation should remain bounded and practical rather than fully theoretical. Fourth, although the core metrics and efficiency profiling are complete for the compared baselines, the full artifact-diagnostic pack has only been re-run for a subset of methods, and user-study data has not yet been collected. Finally, high-resolution generalization beyond the present 256-based protocol remains under-validated. Future work should evaluate richer domain prototype priors, controlled integration horizons, higher-resolution extrapolation, and human preference studies.

Conclusion

We presented OT-coupled latent flow matching for efficient domain-level artistic transfer. The method combines three complementary objectives: a flow-matching loss supervises a time-conditioned velocity field along stochastic latent bridges; terminal sliced-Wasserstein matching pulls the integrated endpoint toward the target style distribution; and kinetic regularization constrains path displacement to preserve content. Under a strict 750-output evaluation, the

resulting 3.9M-parameter model achieves competitive CLIP-style (0.716), the lowest LPIPS among high-style methods (0.451), the best EC score (0.393), and substantially lower training cost than heavier feed-forward baselines—a Pareto-efficient operating point for domain-level artistic transfer.

References

- [1] Mikołaj Bińkowski, Dougal J. Sutherland, Michael Arbel, and Arthur Gretton. Demystifying mmd gans. In *International Conference on Learning Representations*, 2018.
- [2] Jiwoo Chung, Sangeek Hyun, and Jae-Pil Heo. Style injection in diffusion: A training-free approach for adapting large-scale diffusion models for style transfer. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024.
- [3] Marco Cuturi. Sinkhorn distances: Lightspeed computation of optimal transport. *Advances in Neural Information Processing Systems*, 2013.
- [4] Yingying Deng, Fan Tang, Weiming Dong, Wen Sun, Feiyue Huang, and Changsheng Xu. Stytr2: Image style transfer with transformers. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2022.
- [5] Keyan Ding, Kede Ma, Shiqi Wang, and Eero P. Simoncelli. Image quality assessment: Unifying structure and texture similarity. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2020.
- [6] Leon A. Gatys, Alexander S. Ecker, and Matthias Bethge. Image style transfer using convolutional neural networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016.
- [7] Heeseung Gim, Jonghwan Yoo, and Youngjung Uh. Content-adaptive style transfer: A training-free approach with vq autoencoders. In *Proceedings of the Asian Conference on Computer Vision*, 2024.
- [8] Martin Heusel, Hubert Ramsauer, Thomas Unterthiner, Bernhard Nessler, and Sepp Hochreiter. Gans trained by a two time-scale update rule converge to a local nash equilibrium. In *Advances in Neural Information Processing Systems*, 2017.
- [9] Kibeom Hong, Sangryul Jeon, Hyesong Yang, Jianlong Fu, and Hyeran Byun. Aespa-net: Aesthetic pattern-aware style transfer networks. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2023.
- [10] Xuan Huang, Zhen Zhang, Jing Zhang, et al. Aesfa: An aesthetic feature-aware arbitrary neural style transfer. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2024.
- [11] Xun Huang and Serge Belongie. Arbitrary style transfer in real-time with adaptive instance nor-

- malization. In Proceedings of the IEEE International Conference on Computer Vision, 2017.
- [12] Junjie Ke, Qifei Wang, Yilin Wang, Peyman Milanfar, and Feng Yang. Musiq: Multi-scale image quality transformer. Proceedings of the IEEE/CVF International Conference on Computer Vision, 2021.
 - [13] Christian Léonard. A survey of the schrödinger problem and some of its connections with optimal transport. Discrete and Continuous Dynamical Systems - A, 2014.
 - [14] Yahui Liu et al. Pluggable style representation learning for multi-style transfer. In Proceedings of the Asian Conference on Computer Vision, 2024.
 - [15] Dae Young Park and Kwang Hee Lee. Arbitrary style transfer with style-attentional networks. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, 2019.
 - [16] Gabriel Peyré and Marco Cuturi. Computational optimal transport. Foundations and Trends in Machine Learning, 2019.
 - [17] Matthias Wright and Björn Ommer. Artfid: Quantitative evaluation of neural style transfer. arXiv preprint arXiv:2207.12280, 2022.
 - [18] Wei Xia et al. S2wat: Image style transfer via hierarchical vision transformer using strips window attention. In Proceedings of the AAAI Conference on Artificial Intelligence, 2024.
 - [19] Sidi Yang, Tianhe Wu, Shuwei Shi, Shanshan Lao, Yuan Gong, Mingdeng Cao, Jiahao Wang, and Yujiu Yang. Maniq: Multi-dimension attention network for no-reference image quality assessment. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops, 2022.
 - [20] Richard Zhang, Phillip Isola, Alexei A. Efros, Eli Shechtman, and Oliver Wang. The unreasonable effectiveness of deep features as a perceptual metric. Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition, 2018.

Reproducibility Checklist

This paper

- Includes a conceptual outline and/or pseudocode description of AI methods introduced: Yes.
- Clearly delineates statements that are opinions, hypothesis, and speculation from objective facts and results: Yes.
- Provides well marked pedagogical references for less-familiar readers to gain background necessary to replicate the paper: Yes.

Does this paper make theoretical contributions? No. The paper uses a bounded bridge-inspired motivation rather than theorem-level theoretical claims.

Does this paper rely on one or more datasets? Yes.

- A motivation is given for why the experiments are conducted on the selected datasets: Yes.
- All novel datasets introduced in this paper are included in a data appendix: NA.
- All novel datasets introduced in this paper will be made publicly available upon publication of the paper with a license that allows free usage for research purposes: NA.
- All datasets drawn from the existing literature are accompanied by appropriate citations: Yes.
- All datasets drawn from the existing literature are publicly available: Yes.
- All datasets that are not publicly available are described in detail, with explanation why publicly available alternatives are not scientifically satisfying: NA.

Does this paper include computational experiments? Yes.

- This paper states the number and range of values tried per (hyper-)parameter during development of the paper, along with the criterion used for selecting the final parameter setting: Partial.
- Any code required for pre-processing data is included in the appendix: No.
- All source code required for conducting and analyzing the experiments is included in a code appendix: No.
- All source code required for conducting and analyzing the experiments will be made publicly available upon publication of the paper with a license that allows free usage for research purposes: Partial.
- All source code implementing new methods have comments detailing the implementation, with references to the paper where each step comes from: Partial.
- If an algorithm depends on randomness, then the method used for setting seeds is described in a way sufficient to allow replication of results: Partial.
- This paper specifies the computing infrastructure used for running experiments, including hardware/software details: Partial.
- This paper formally describes evaluation metrics used and explains the motivation for choosing these metrics: Yes.
- This paper states the number of algorithm runs used to compute each reported result: Yes.
- Analysis of experiments goes beyond single-dimensional summaries of performance to include measures of variation, confidence, or other distributional information: Yes.

- The significance of any improvement or decrease in performance is judged using appropriate statistical tests: No.
- This paper lists all final (hyper-)parameters used for each model/algorithm in the paper's experiments: Partial.